



MORE SQL: COMPLEX QUERIES, TRIGGERS, VIEWS, AND SCHEMA MODIFICATION

Chapter 7

Chapter Outline

- More Complex SQL Retrieval Queries
- Specifying Semantic Constraints as Assertions and Actions as Triggers
- Views (Virtual Tables) in SQL
- Schema Modification in SQL

MORE COMPLEX SQL RETRIEVAL QUERIES

Chapter 7.1

More Complex SQL Retrieval Queries

- Additional feature allow users to specify more complex and interesting retrievals from database.
 - Such features:
 - Nested queries (중첩 질의)
 - Joined tables (Natural Join)
 - *Outer* joins in the `FROM` clause
 - Views (Derived Tables), Assertions, Triggers
 - Aggregate functions (집계 함수)
 - Grouping
- This chapter focuses on learning what they are and how to use them in SQL.

Comparisons Involving NULL and Three-Valued Logic

- SQL uses a *three-valued* logic.
 - The result of evaluating an expression falls in:
 - TRUE, FALSE, and UNKNOWN
 - NULL = NULL cannot be evaluated.
- Logical connectives (truth table) in the three-Valued Logic

AND	TRUE	FALSE	UNKNOWN
TRUE	TRUE	FALSE	UNKNOWN
FALSE	FALSE	FALSE	FALSE
UNKNOWN	UNKNOWN	FALSE	UNKNOWN
OR	TRUE	FALSE	UNKNOWN
TRUE	TRUE	TRUE	TRUE
FALSE	TRUE	FALSE	UNKNOWN
UNKNOWN	TRUE	UNKNOWN	UNKNOWN
NOT	TRUE	FALSE	UNKNOWN
TRUE	FALSE		
FALSE	TRUE		
UNKNOWN	UNKNOWN		

Comparisons Involving NULL and Three-Valued Logic (Cont'd)

- SQL allows queries checking whether an attribute value **IS** NULL.

```
SELECT Fname, Lname  
FROM EMPLOYEE  
WHERE Super_ssn IS NULL;
```

- The above query looks for the names of all employees who do not have supervisors.

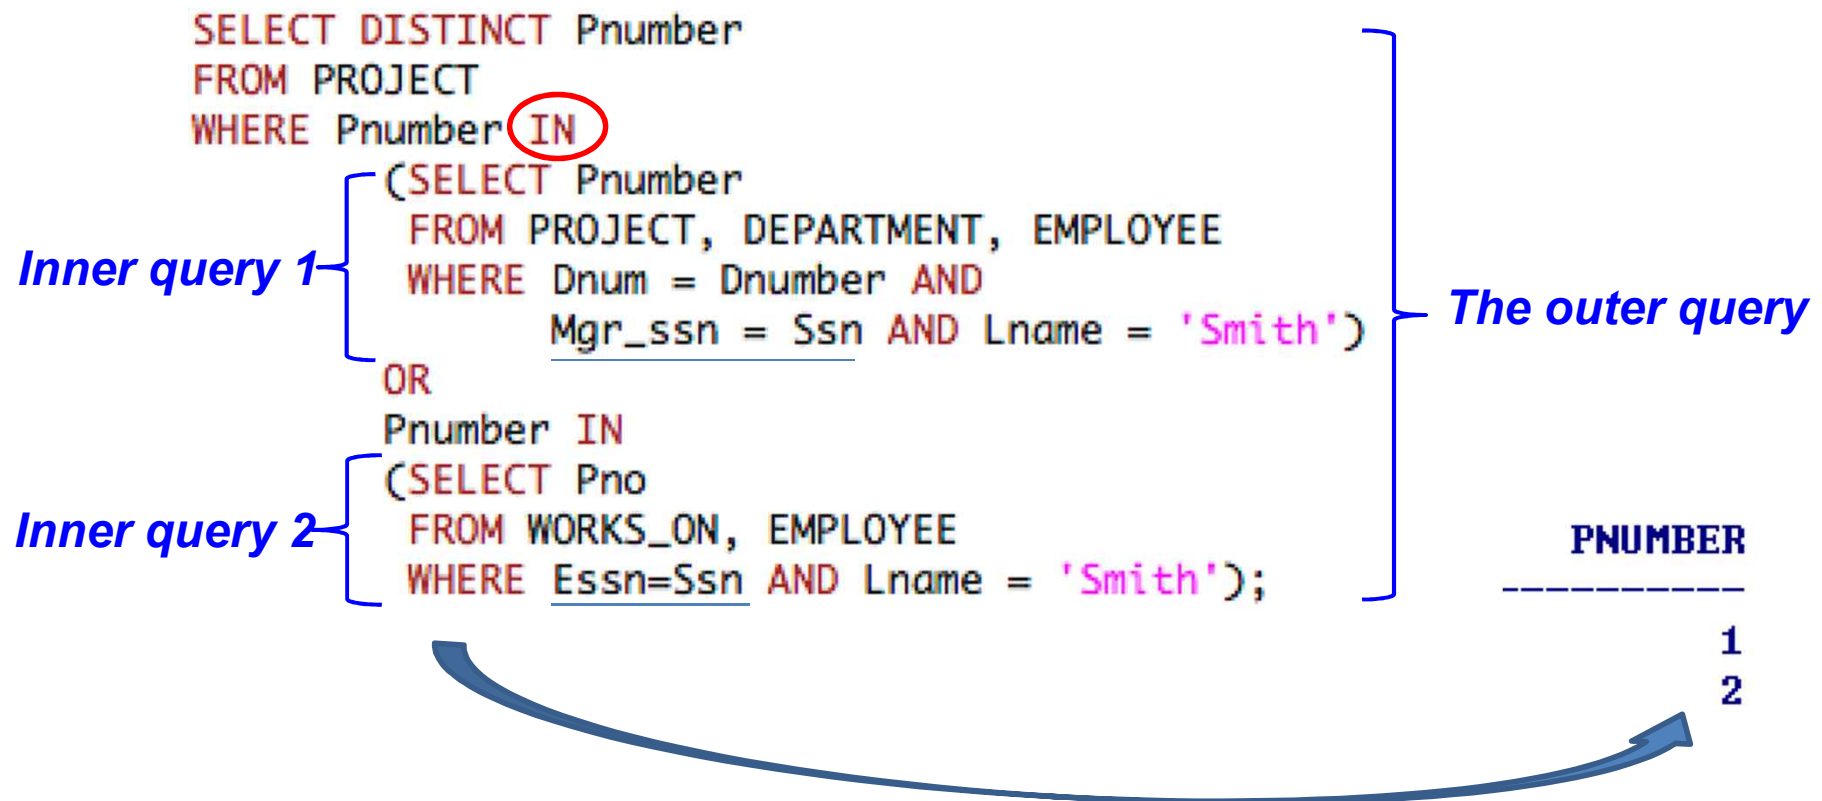


<u>FNAME</u>	<u>LNAME</u>
James	Borg

- For this result, we assume that we restore the COMPANY database before the update and deletion done in Lab #5-4.

Nested Queries

- Have complete a *select-from-where* block(s), called a (nested) **subquery** or an **inner query**, within WHERE clause of another query, called an **outer query**



(Nested Queries (Cont'd))

Set/Multiset Comparison Operator: **IN**

- Compares value v with a set (or multiset) of values V
- Evaluates to TRUE if v is one of the elements in V
- Can be used for comparing “tuples of values”
 - To do the comparison, place the tuples within ‘ () ’.

```
SELECT DISTINCT Essn
FROM WORKS_ON
WHERE (Pno, Hours) IN (SELECT Pno, Hours
                       FROM WORKS_ON
                       WHERE Essn = '123456789');
```




ESSN
123456789

Nested Queries (Cont'd)

- We can use other comparison operator to compare a single value *v*.
- **ALL**: value must exceed “all” values from nested query

```
SELECT Lname, Fname
FROM   EMPLOYEE
WHERE  Salary > ALL ( SELECT Salary
                      FROM   EMPLOYEE
                      WHERE   Dno = 5);
```




LNAME	FNAME
Wallace	Jennifer
Borg	James

Nested Queries (Cont'd)

- To avoid potential errors and ambiguities, create “aliases” for all tables referenced in an SQL query
- Example) “Retrieve the name of each employee who
 - 1) Has a dependent with the same first name, and
 - 2) Is the same gender as that of the employee.

```
SELECT  E.Fname, E.Lname
FROM    EMPLOYEE E
WHERE   E.Ssn IN ( SELECT      Essn
                   FROM        DEPENDENT D
                   WHERE E.Fname = D.Dependent_name
                   AND E.Sex   = D.Sex );
```




no rows selected

- Called a **correlated nested query** (상호 연관된 중첩 질의)
 - Evaluated once for each tuple in the outer query

Nested Queries (Cont'd)

- Nested queries using the '=' or '**IN**' can be rewritten into one single query with a join condition.
- Ex) The previous query can be written in the following:

```
SELECT  E.Fname, E.Lname
FROM    EMPLOYEE E, DEPENDENT D
WHERE   E.Ssn = D.Essn
        AND E.Sex = D.Sex
        AND E.Fname = D.Dependent_name;
```



no rows selected

The (NOT) EXISTS Functions in SQL for Correlating Queries

- (NOT) EXISTS function
 - Check whether the result of a correlated nested query is 'empty' or not.
 - Can be used in conjunction with a correlated nested query
 - Is a Boolean function that returns a TRUE or FALSE result.
 - If there is no tuple returned by the correlated nested query, then EXISTS (NOT EXISTS) returns TRUE (or FALSE).

```
SELECT Fname, Lname
FROM Employee
WHERE EXISTS (SELECT *
              FROM DEPENDENT
              WHERE Ssn = Essn)
AND EXISTS (SELECT *
            FROM Department
            WHERE Ssn= Mgr_Ssn);
```

FNAME	LNAME
Franklin	Wong
Jennifer	Wallace



Use of **NOT EXISTS**

- To achieve the “for all” (\forall), a universal quantifier (정량자: “모든 것에 대해”), effect, we may use double negation this way in SQL:
(...하지 않은 튜플들은 존재하지 않는다 -> ...한 튜플들만 존재한다.)
- “Retrieve the name of employees working on “ALL” projects controlled by Dno = 5.”

=> List names of those employees for whom there does **NOT** exist a project controlled by department no 5 that they do **NOT** work on.

(직원들이 일도 하지 않는 5번 부서에서 관리되는 과제가 없는 그러한 직원들의 이름을 나열하라.)

```
SELECT  E.Fname, E.Lname
FROM    EMPLOYEE E
WHERE   NOT EXISTS ( (SELECT  P.Pnumber
                     FROM    PROJECT P
                     WHERE   Dnum = 5)
               MINUS  -- 'EXCEPT' in the SQL standard
               (SELECT  W.Pno
                FROM    WORKS_ON W
                WHERE   E.Ssn = w.Essn));
```

no rows selected

What does this mean?

Appendix: Use of **NOT EXISTS** (Cont'd)

- The previous query can be rewritten in a more complex way of using two-level nesting:

```
SELECT E.Fname, E.Lname
FROM   EMPLOYEE E
WHERE  NOT EXISTS (SELECT
                  FROM
                  WHERE
```

* 5번 부서에서 관리하는 과제 중에,
한 과제라도 빠져 있으면 그 직원은
본 질의의 결과에 포함될 수 없음

(e.g., E.Ssn = '123456789')

(3번 과제 참여x)

or '66688444' (1,2번 과제 x)

or '453453453' (3번 과제 x)

or '33445555' (1번 과제 x).


E.Ssn = '123456789' 의 경우,

666884444	3	40.0
-----------	---	------

(in WORKS_ON)
가 존재하게 하게 되어, 선택받지 못함

```
*
WORKS_ON B
(B.Pno IN (SELECT P.Pnumber
            FROM   PROJECT P
            WHERE  P.Dnum = 5)
AND
NOT EXISTS (SELECT *
            FROM WORKS_ON C
            WHERE C.Essn = E.ssn
            AND C.Pno = B.Pno)));
```

no rows selected



“Select each employee, such that there does **NOT EXIST** a project controlled by department 5 that the employee does **NOT WORK ON**.”

SQL Function: **UNIQUE** (Q)

- Returns
 - TRUE if there are no duplicate tuples in the result query Q
 - FALSE, otherwise.
- Can be used to test whether the result of a nested query is a set (**no duplicates**) or a multiset (**duplicates exist**).

Explicit Sets and Renaming in SQL

- An *explicit* set of values can be used in the WHERE clause.

```
SELECT DISTINCT Essn
FROM   WORKS_ON
WHERE  Pno IN (1, 2, 3);
```

ESSN

333445555
453453453
123456789
666884444

- Attribute renaming: use “AS” followed by whatever name is legal. (Discussed last time)

```
SELECT E.Lname AS Employee_name,
       S.Lname As Supervisor_name
FROM   EMPLOYEE E, EMPLOYEE S
WHERE  E.Super_ssn = S.Ssn;
```

EMPLOYEE_NAME	SUPERVISOR_NAME
Wallace	Borg
Wong	Borg
Zelaya	Wallace
Jabbar	Wallace
Smith	Wong
Narayan	Wong
English	Wong


7 rows selected.

Joined Tables in SQL and Inner Joins

- **Joined Tables**

- Concept: Users can be permitted to specify a table resulting from a join operation *in the FROM clause* of a query.

```
SELECT Fname, Lname, Address
FROM   (EMPLOYEE JOIN DEPARTMENT ON Dno = Dnumber)
WHERE  Dname = 'Research';
```



FNAME	LNAME	ADDRESS
Franklin	Wong	638 Uoss, Houston, TX
John	Smith	731 Fondren, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX


- Contains a single joined table; Such join may also be called **inner join** (to be discussed later) (working for only matching tuples).

Different Types of JOINed Tables in SQL

- Users can specify *different* types of join:
 - **NATURAL JOIN** (most representative of an inner join)
 - For R (left table) \bowtie S (right table), *no* join condition is specified.
 - Same as creating an implicit **EQUIJOIN** condition for *each pair of attributes with the same name* from R and S

```
-- For renaming as Dno
CREATE TABLE DEPT AS
    SELECT Dname, Dnumber as Dno, Mgr_ssn, Mgr_start_date
    FROM DEPARTMENT;
-- Natural join
SELECT Fname, Lname, Address
FROM EMPLOYEE NATURAL JOIN DEPT
WHERE Dname = 'Research';
```

You should rename attributes of one relation so it can be joined with another using NATURAL JOIN.



FNAME	LNAME	ADDRESS
Franklin	Wong	638 Uoss, Houston, TX
John	Smith	731 Fondren, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

Different Types of JOINed Tables in SQL (Cont'd)

- Users can specify different types of join (Cont'd):
 - **INNER JOIN** (vs. OUTER JOIN)
 - Default type of join in a joined table
 - Tuple is included in the result *only* if a matching tuple exists in the other relation. What if we'd like to see the result including non-matching tuples?
 - **LEFT (RIGHT) OUTER JOIN**
 - **EVERY** tuple in left (right) table (say, R (S)) must appear in the result.
 - If no matching tuple,
 - Padded with `NULL` values for attributes of right table (say, S (R))

- Example

```
SELECT E.Lname AS Supervisee_Name,  
       S.Lname AS Supervisor_Name  
FROM Employee E LEFT OUTER JOIN EMPLOYEE S  
ON E.Super_ssn = S.Ssn
```


SUPERVISEE_NAME	SUPERVISOR_NAME
Wallace	Borg
Wong	Borg
Zelaya	Wallace
Jabbar	Wallace
Smith	Wong
Narayan	Wong
English	Wong
Borg	

Multiway JOIN in the FROM clause

- **FULL OUTER JOIN**: combines the result if LEFT and RIGHT OUTER JOIN
- A “multiway” join can be specified by nesting JOIN specifications.
 - Example)

```
SELECT Pnumber, Dnum, Lname, Address, Bdate  
FROM ((PROJECT JOIN DEPARTMENT ON Dnum=Dnumber) JOIN EMPLOYEE ON Mgr_ssn=Ssn)  
WHERE Plocation='Stafford';
```

PNUMBER	DNUM	LNAME	ADDRESS	BDATE
10	4	Wallace	291 Berry, Bellaire, TX	20-JUN-41
30	4	Wallace	291 Berry, Bellaire, TX	20-JUN-41



Aggregate Functions in SQL

- Why? To summarize information from multiple tuples into a **single-tuple** summary
- Built-in aggregate functions: COUNT, SUM, MAX, MIN, AVG
- Typically, grouping via GROUP BY clause
 - Create subgroups of tuples before summarizing
- To select (or, apply condition to) entire groups, HAVING clause is used.
- Aggregate functions can be used in the SELECT clause and a HAVING clause.

Aggregations Applied for Entire Tuples

```
SELECT SUM (Salary), MAX (Salary), MIN (Salary), AVG (Salary)
FROM EMPLOYEE;
```

SUM<SALARY>	MAX<SALARY>	MIN<SALARY>	AVG<SALARY>
281000	55000	25000	35125

```
SELECT COUNT(*) as NumEmps, SUM (Salary) AS Total_Sal, MAX (Salary) AS Highest_Sal,
MIN (Salary) AS Lowest_Sal, AVG (Salary) AS Average_Sal
FROM EMPLOYEE, DEPARTMENT
WHERE Dno = Dnumber AND Dname = 'Research';
```

NUMEMPS	TOTAL_SAL	HIGHEST_SAL	LOWEST_SAL	AVERAGE_SAL
4	133000	40000	25000	33250

```
SELECT COUNT (E.Lname)
FROM Employee E LEFT OUTER JOIN EMPLOYEE S
ON E.Super_ssn = S.Ssn;
```

COUNT<E.LNAME>
8

```
SELECT COUNT (S.Lname)
FROM Employee E LEFT OUTER JOIN EMPLOYEE S
ON E.Super_ssn = S.Ssn;
```

COUNT<S.LNAME>
7

```
SELECT COUNT (*)
FROM Employee E LEFT OUTER JOIN EMPLOYEE S
ON E.Super_ssn = S.Ssn;
```

COUNT<*>
8

NULLs are discarded for counting values, EXCEPT for **tuples**.

Grouping: The **GROUP BY** Clause

- **Partition** relation into subsets of tuples
 - Based on *grouping attributes*: having the same value for them
 - Apply function to each such group independently

- **GROUP BY clause**

- Specifies grouping attributes
- The grouping attribute **MUST** appear in the **SELECT** clause.

```
SELECT Dno, COUNT (*) as nEmps, AVG (Salary) as avgSal
FROM EMPLOYEE
GROUP BY Dno;
```



DNO	NEMPS	AUGSAL
1	1	55000
4	3	31000
5	4	33250

Fname	Minit	Lname	Ssn	...	Salary	Super_ssn	Dno
John	B	Smith	123456789		30000	333445555	5
Franklin	T	Wong	333445555		40000	888665555	5
Ramesh	K	Narayan	666884444		38000	333445555	5
Joyce	A	English	453453453	...	25000	333445555	5
Alicia	J	Zelaya	999887777		25000	987654321	4
Jennifer	S	Wallace	987654321		43000	888665555	4
Ahmad	V	Jabbar	987987987		25000	987654321	4
James	E	Bong	888665555		55000	NULL	1

Dno	Count (*)	Avg (Salary)
5	4	33250
4	3	31000
1	1	55000

Result of Q24

Grouping EMPLOYEE tuples by the value of Dno

Grouping: The **GROUP BY** Clause

- GROUP BY (GB) clause (Cont'd)
 - Can be applied to the result of JOIN

```
SELECT Pnumber, Pname, COUNT (*) as numEmps
FROM   PROJECT, WORKS_ON
WHERE  Pnumber=Pno
GROUP BY Pnumber, Pname
ORDER BY Pnumber, Pname;
```

GB *applied to the result of JOIN and then sorted*



PNUMBER	PNAME	NUMEMPS
1	ProductX	2
2	ProductY	3
3	ProductZ	2
10	Computerization	3
20	Reorganization	3
30	NewBenefits	3

Grouping: The **GROUP BY** Clause with **HAVING** Clause

- HAVING clause
 - Provides a condition to select or reject an entire group

Q26: `SELECT Pnumber, Pname, COUNT (*) as numEmps
FROM PROJECT, WORKS_ON
WHERE Pnumber=Pno
GROUP BY Pnumber, Pname
HAVING COUNT(*) > 2;`

Pname	Pnumber	...	Essn	Pno	Hours
ProductX	1		123456789	1	32.5
ProductX	1		453453453	1	20.0
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
ProductZ	3		666884444	3	40.0
ProductZ	3		333445555	3	10.0
Computerization	10	...	333445555	10	10.0
Computerization	10		999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

These groups are not selected by the HAVING condition of Q26.

After applying the WHERE clause but before applying HAVING

Grouping: The **GROUP BY** Clause with **HAVING** Clause (Cont'd)

- HAVING clause (Cont'd)
 - Provides a condition to select or reject an entire group

Q26: `SELECT Pnumber, Pname, COUNT (*) as numEmps
FROM PROJECT, WORKS_ON
WHERE Pnumber=Pno
GROUP BY Pnumber, Pname
HAVING COUNT(*) > 2;`

Pname	Pnumber	...	Essn	Pno	Hours
ProductY	2		123456789	2	7.5
ProductY	2		453453453	2	20.0
ProductY	2		333445555	2	10.0
Computerization	10		333445555	10	10.0
Computerization	10	...	999887777	10	10.0
Computerization	10		987987987	10	35.0
Reorganization	20		333445555	20	10.0
Reorganization	20		987654321	20	15.0
Reorganization	20		888665555	20	NULL
Newbenefits	30		987987987	30	5.0
Newbenefits	30		987654321	30	20.0
Newbenefits	30		999887777	30	30.0

Pname	Count (*)
ProductY	3
Computerization	3
Reorganization	3
Newbenefits	3

Result of Q26
(Pnumber not shown)

PNUMBER	PNAME	NUMEMPS
20	Reorganization	3
30	NewBenefits	3
10	Computerization	3
2	ProductY	3

After applying the HAVING clause condition

Grouping: The **GROUP BY** Clause with **HAVING** Clause (Cont'd)

- HAVING clause (Cont'd)
 - Provides a condition to select or reject an entire group

Q28: “For each department with **≥ 2** employees, retrieve the *department number* and the *number of its employee* who’re earning **$> \$40,000$** .”

```
SELECT Dnumber, COUNT(*) as nEmps
FROM   DEPARTMENT, EMPLOYEE
WHERE  Dnumber = Dno AND Salary > 40000
      AND Dnumber IN
      (SELECT Dno
       FROM   EMPLOYEE
       GROUP BY Dno
       HAVING COUNT(*) >= 2)
GROUP BY Dnumber;
```



DNUMBER	NEMPS
4	1

WITH Clause

- Allows a user to define a (temporary) table that will only be used in a “particular” query.
- Used for convenience to create a temporary “View” and use that immediately in a query: called an *in-line view*
- Q28 can be rewritten with the WITH clause:

```
WITH SOMEDEPTS AS
  (SELECT Dno
   FROM EMPLOYEE
   GROUP BY Dno
   HAVING COUNT (*) >= 2)
SELECT e.Dno, COUNT (*) AS nEmps
FROM EMPLOYEE e, SOMEDEPTS b
WHERE Salary > 40000 AND e.Dno = b.Dno
GROUP BY e.Dno;
```



DNO	NEMPS
4	1

Use of **CASE** Clause

- Used when a value can be different based on “certain conditions”
- Can be used in *any* part of an SQL query where a value is expected
- Applicable when *querying, inserting, or updating* tuples

```
UPDATE  EMPLOYEE
SET Salary =
CASE    WHEN    Dno = 5 THEN    Salary + 2000
        WHEN    Dno = 4 THEN    Salary + 1500
        WHEN    Dno = 1 THEN    Salary + 3000
```

Recursive Queries in SQL

- Can be used to keep track of the relationship between tuples of the same type: e.g., employee vs. supervisor
 - Such relationship is described by the FK, `Super_ssn` of `EMPLOYEE`.

```
WITH RECURSIVE SUP_EMP (SupSsn,
EmpSsn) AS
(SELECT Super_Ssn, Ssn
FROM EMPLOYEE
UNION
SELECT E.Ssn, S.SupSsn
FROM EMPLOYEE E, SUP_EMP S
WHERE E.Super_Ssn = S.EmpSsn)
SELECT *
FROM SUP_EMP;
```

[In the SQL standard]

```
SELECT Super_Ssn, Ssn as EmpSsn
FROM EMPLOYEE
START WITH Ssn = '123456789'
CONNECT BY PRIOR Super_Ssn = Ssn;
```

[In Oracle]



SUPER_SSN	EMPSSN
333445555	123456789
888665555	333445555
	888665555

Reminder: EXPANDED Block Structure of SQL Queries

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```

SPECIFYING CONSTRAINTS AS **ASSERTIONS** AND ACTIONS AS **TRIGGERS**

Chapter 7.2

CREATE ASSERTION

- Can allow users to specify “general” constraints via **declarative assertions**.
 - The constraints do not fall into any of the categories of key (or unique), entity, not-null, referential integrity constraints.
- Specifies a query that selects any tuples violating the desired condition (set by the users).
- Use *only* in cases that cannot be specified by a simple CHECK which applies to individual attributes and domains.

```
CREATE ASSERTION SALARY_CONSTRAINT  
CHECK (NOT EXISTS (SELECT *  
                    FROM EMPLOYEE E, EMPLOYEE M,  
                        DEPARTMENT D  
                    WHERE E.Salary > M.Salary  
                          AND E.Dno = D.Dnumber      (Not implemented  
                          AND D.Mgr_ssn = M.Ssn) );    by Oracle)
```

Triggers (트리거)

- Convenient to specify the *type of action* to be taken when *certain events occur* and *certain conditions are satisfied*
 - “If an employee exceeds a travel expense limit, notify his manager.”
 - Used to monitor the database
- Typical trigger has three components:
 - **Event(s), Condition, Action (ECA)**
 - So the trigger is regarded as an ECA rule.
 - These make it a rule for an “active” database, which is out of scope in this course.
 - For those who are further interested, refer to Section 26.1.

Triggers – How to Use? (Cont'd)

```
CREATE OR REPLACE  
TRIGGER SALARY_VIOLATION  
-- Event
```

- Inserting a new employee record
- Changing an employee's salary
- Changing an employee's supervisor

```
BEFORE INSERT OR UPDATE OF Salary, Supervisor_ssn ON EMPLOYEE  
FOR EACH ROW
```

-- Condition: Determines whether the rule action should be executed

```
WHEN (:NEW.SALARY > (SELECT Salary  
FROM EMPLOYEE  
WHERE Ssn = NEW.Supervisor_Ssn))
```

-- Action: usually a sequence of SQL statements, a transaction, or PSM
INFORM_SUPERVISOR (NEW.Supervisor.Ssn, NEW.Ssn);

- Means that the trigger should be executed "BEFORE" the triggering operation is executed.

Called a **stored procedure** (저장 프로시저): a program module stored by the DBMS at the database server; in the SQL standard, called **persistent stored modules (PSM)** (영속 저장 모듈)

Not executed in Oracle as it is....

Triggers – Another Example (Cont'd)

```
-- Declaration of a trigger
CREATE OR REPLACE TRIGGER knu.SALARY_VIOLATION
BEFORE INSERT OR UPDATE ON knu.EMPLOYEE
FOR EACH ROW
WHEN (NEW.SALARY > 100000)
DECLARE
    sal_diff number;
BEGIN
    sal_diff := :NEW.salary - :OLD.salary;
    dbms_output.put_line('Old salary: ' || :OLD.salary);
    dbms_output.put_line('New salary: ' || :NEW.salary);
    dbms_output.put_line('Salary difference: ' || sal_diff);
END;
```

A Block of
**Oracle
PL/SQL**

Execute {

One of the functions in the package

A system package provided by Oracle

“Report when a employee’s new salary exceeds a salary cap, \$100K.”

Triggers

– On Oracle

(Cont'd)

```
SQL> CREATE OR REPLACE TRIGGER knu.SALARY_VIOLATION
 2  BEFORE INSERT OR UPDATE ON knu.EMPLOYEE
 3  FOR EACH ROW
 4  WHEN (NEW.SALARY > 100000)
 5  DECLARE
 6      sal_diff number;
 7  BEGIN
 8      sal_diff := :NEW.salary - :OLD.salary;
 9      dbms_output.put_line('Old salary: ' || :OLD.salary);
10      dbms_output.put_line('New salary: ' || :NEW.salary);
11      dbms_output.put_line('Salary difference: ' || sal_diff);
12  END;
13  /
```

Trigger created.

```
SQL> ALTER TRIGGER SALARY_VIOLATION ENABLE;
-- Enable trigger manually
Trigger altered.
```

```
SQL>
SQL> SET SERVEROUTPUT ON -- Enable print screen
SQL>
```

```
SQL> UPDATE knu.EMPLOYEE
 2  SET Salary = Salary*2
 3  WHERE Ssn = '888665555';
```

Old salary: 55000

New salary: 110000

Salary difference: 55000

1 row updated.

```
SQL>
```

```
SQL> DROP TRIGGER knu.SALARY_VIOLATION;
-- Drop trigger
```

Trigger dropped.

VIEWS(VIRTUAL TABLES) IN SQL

Chapter 7.3

Views (Virtual Tables) in SQL

- Concept of a view in SQL
 - Single table **derived from** other tables
 - Somewhat different than “user view” involving many relations
 - Considered to be a *virtual* table that is not “necessarily populated”
- A view can be thought of as a way of specifying a table that needs to be referenced frequently (though it may not exist physically)
 - Often, the view is used for the purpose of caching the result of joins that are “frequently requested” (join cost vs. **space** (compromised))
 - Ex) “Retrieve the **employee** name and the **project** names that the employee **works on**.”
 - Joins `EMPLOYEE`, `WORKS_ON`, and `PROJECT` every time this query issued.
 - If a view is defined on these joins, then it works for “single-table” retrievals.

Specification of Views in SQL: **CREATE VIEW**

- Give view (virtual table) name, list attribute names, and include view definition—a “query” to specify the contents of the view

```
V1: CREATE VIEW WORKS_ON1 AS
      SELECT Fname, Lname, Pname, Hours
      FROM   EMPLOYEE, PROJECT, WORKS_ON
      WHERE  Ssn = Essn AND Pno = Pnumber;
```

View created.

```
SQL> desc  WORKS_ON1
```

Name	Null?	Type
FNAME	NOT NULL	VARCHAR2(15)
LNAME		VARCHAR2(15)
PNAME	NOT NULL	VARCHAR2(15)
HOURS		NUMBER(3,1)

- V1 *inherits* the names of view attributes from the defining tables.

Specification of Views in SQL: **CREATE VIEW** (Cont'd)

- Give view (virtual table) name, list attribute names, and include view definition—a “query” to specify the contents of the view

```
V2: CREATE VIEW DEPT_INFO (Dept_name, No_of_emps, Total_sal) AS
      SELECT    Dname, COUNT(*), SUM(Salary)
      FROM      DEPARTMENT, EMPLOYEE
      WHERE     Dnumber = Dno
      GROUP BY  Dname;
```

View created.

```
SQL> desc DEPT_INFO;
```

Name	Null?	Type
DEPT_NAME	NOT NULL	VARCHAR2(15)
NO_OF_EMPS		NUMBER
TOTAL_SAL		NUMBER

- V2 explicitly specifies new attribute names for the view by a one-to-one correspondence.

Specification of Views in SQL (Cont'd)

- Once a View is defined, SQL queries can use the View relation in the FROM clause.
 - Ex) Accessing the defined view: WORKS_ON1

```
SELECT Fname, Lname  
FROM   WORKS_ON1  
WHERE  Pname = 'ProductX';
```



FNAME	LNAME
John	Smith
Joyce	English

Specification of Views in SQL (Cont'd)

- Why using a view? Advantages of defining a view
 - 1) Simplification of the specification of certain queries.
 - 2) Provision of a security and authorization mechanism
 - 3) Saving (multiple) expensive join cost by space *if materialized*
- **DROP VIEW**
 - Dispose of a view.

```
SQL> drop view WORKS_ON1;  
  
View dropped.  
  
SQL>
```

View Implementation

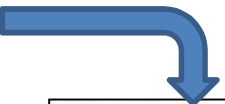
- A view is supposed to be **always up-to-date**. Why?
 - If we *modify* tuples in the base relation on which the view is defined, then the view must “automatically” reflect these changes.
 - The view should have to be **realized** or **materialized**
 - At the time of *specifying a query on the view* but not of defining the view
- The DBMS is responsible for keeping the view up-to-date
 - **NOT** the user to make sure that the view is newest.
 - Question: Then how does the DBMS let the view to be up-to-date?
 - Problem: not easy to “efficiently” implement a view for querying

View Implementation (Cont'd)

- Strategy 1) **Query modification** approach
 - Compute the view as and when needed.
 - Let's not store the view permanently!
 - Modify view query into a query on underlying base tables.

`SELECT Fname, Lname
FROM WORKS_ON1 -- View V1
WHERE Pname = 'ProductX';`

transform



```
SELECT Fname, Lname  
FROM EMPLOYEE, PROJECT, WORKS_ON  
WHERE Ssn = Essn AND Pno = Pnumber  
AND Pname = 'ProductX';
```

- Any *problem*??
 - **Inefficient** for views defined via “complex” queries, which take so long, or *time-consuming*, to execute

View Implementation (Cont'd)

- Strategy 2) **View materialization** approach
 - *Physically create a temporary view table* when the view is first queried.
 - And keep that table on the assumption that other queries on the view will come later.
 - Requires efficient strategy for automatically updating the view table when the base tables are updated.
 - Incremental update strategy for materialized views
 - Means the DBMS can determine what new tuples must be inserted, deleted, or modified in a *materialized view table*
 - When a database update is applied *to one of the (defining) base tables*.
 - A materialized table is maintained as long as it is being queried.
 - If no query on the view for a certain period of time, then the table is automatically removed. Later recomputed from scratch, if accessed again.


View Implementation (Cont'd)

- Multiple ways to implement materialization:
 - *Immediate update* (c.f., “write-through”)
 - Updates a view *as soon as the base tables are changed*.
 - *Lazy update* (c.f., “write-back”)
 - Updates a view whenever a view query requests
 - *Periodic update*
 - Updates the view periodically
- Note that in the *latter* strategy, a view query may not get an up-to-date result.
 - This is commonly used in Population Survey, Monthly Sale Record Retrievals, Banks, Retail store operations, etc.

View Update

- In many cases, modifying a view table via INSERT/DELETE/UPDATE command is not possible. Why? It may mean many...

```
UPDATE WORKS_ON1
SET   Pname = 'ProductY'
WHERE Lname = 'Smith' AND Fname = 'John'
AND   Pname = 'ProductX';
```



```
UPDATE WORKS_ON -- base table
SET   Pno = (SELECT Pnumber
             FROM PROJECT
             WHERE Pname = 'ProductY')
WHERE Essn IN (SELECT Ssn
              FROM   EMPLOYEE
              WHERE  Lname = 'Smith'
                  AND Fname = 'John')
AND Pno = (SELECT Pnumber
          FROM   PROJECT
          WHERE  Pname = 'ProductX');
```

```
UPDATE PROJECT -- base table
SET Pname = 'ProductY'
WHERE Pname = 'ProductX';
```

But Oracle says, “**cannot modify a column which maps to a non key-preserved table**” due to the existing tuple with the same attribute values as what the tuple to be updated has.

View Update (Cont'd)

- An update on a view defined on a *single table* without any *aggregate functions*
 - Can be translated to an update on underlying base table.
- What if there EXISTS such an aggregate function?

```
UPDATE DEPT_INFO -- View V2
SET     Total_sal=100000
WHERE   Dname='Research';
```

*Do you think this update is permitted?
Why? Or Why **NOT**?*

Views as Authorization Mechanism

- Suppose a certain user is only allowed to see employee information for employees that work for department 5.

```
CREATE VIEW      DEPT5EMP  AS
      SELECT    *
      FROM      EMPLOYEE
      WHERE     Dno = 5;
```

- The DBA may grant to that user the privilege to query the view but not the base table `EMPLOYEE` itself.
- This user then won't be able to see other employee tuples when the view is queried, except the information of employees in DEPT #5.
- This way, view can be used to hide certain attributes or tuples from “unauthorized users.”

SCHEMA CHANGE STATEMENTS IN SQL

Chapter 7.4

Schema Evolution Commands

- Can be used to alter a schema by adding or dropping tables/views, attributes, constraints, and other schema constructs
 - Why? DBA may want to change the schema while the database is operational.
- Do not require recompiling the database schema;
 - Convenient, quick.
 - But ensures that the changes do not affect the rest of the database and make it consistent.

The **DROP** Command

- Used to drop named schema elements: tables, domains, or constraints
- DROP behavior options: **CASCADE** and **RESTRICT**
- Example
 - **DROP SCHEMA** COMPANY **CASCADE**; -- DON'T DO THIS
UNLESS COMPLETE SURE
 - Removes the schema and all its elements including tables, views, constraints, etc.
 - **RESTRICT**: Proceeds with the removal only if there's *no* element in it.
 - **DROP TABLE** DEPENDENT **CASCADE**;
 - Removes the relation and its definition from the catalog
 - **RESTRICT**: Proceeds with the removal only if no reference to it

The **ALTER TABLE** Command

- The actions include
 - Adding or dropping a column (attribute)
 - Changing a column definition
 - Adding or dropping table constraints

- Example:

```
ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job VARCHAR(12);
```

```
ALTER TABLE COMPANY.DEPARTMENT DROP COLUMN Address CASCADE;
```

- **CASCADE**: All constraints and views referencing the column (Address) are dropped with it.
- **RESTRICT**: removes only if no views/constraints reference the column

```
ALTER TABLE COMPANY.EMPLOYEE DROP CONSTRAINT EMPSUPERFK CASCADE;
```

Default Values

- Can be dropped and altered:

```
ALTER TABLE COMPANY.DEPARTMENT  
    ALTER COLUMN Mgr_ssn DROP DEFAULT;
```

```
ALTER TABLE COMPANY.DEPARTMENT  
    ALTER COLUMN Mgr_ssn SET DEFAULT '333445555';
```

SUMMARY OF SQL SYNTAX

```

CREATE TABLE <table name> ( <column name> <column type> [ <attribute constraint> ]
                             { , <column name> <column type> [ <attribute constraint> ] }
                             [ <table constraint> { , <table constraint> } ] )

DROP TABLE <table name>
ALTER TABLE <table name> ADD <column name> <column type>

SELECT [ DISTINCT ] <attribute list>
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }
[ WHERE <condition> ]
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]

<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) )
                  { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) } )

<grouping attributes> ::= <column name> { , <column name> }

<order> ::= ( ASC | DESC )

INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) } )
| <select statement>

DELETE FROM <table name>
[ WHERE <selection condition> ]

UPDATE <table name>
SET <column name> = <value expression> { , <column name> = <value expression> }
[ WHERE <selection condition> ]

CREATE [ UNIQUE ] INDEX <index name>
ON <table name> ( <column name> [ <order> ] { , <column name> [ <order> ] } )
[ CLUSTER ]

DROP INDEX <index name>

CREATE VIEW <view name> [ ( <column name> { , <column name> } ) ]
AS <select statement>

DROP VIEW <view name>

```

NOTE: The commands for creating and dropping indexes are not part of standard SQL.