



PHYSICAL DATABASE DESIGN

Chapters 16

- Disk Storage, Basic File Structures

Prof. Young-Kyoon Suh

Chapter Outline

- Physical Database Design
 - Storage Hierarchy
 - Disk Architecture
- Basic File Structure: Primary File Organization
 - **Heap** file
 - **Sorted** file
 - **Hash** file (if time)

Introduction

- Databases typically stored on magnetic disks
 - Database are accessed using physical database file structures.
- The process of physical database design involves choosing the particular data organization techniques.
 - Best fitting the requirements of application among a variety of options.
- In today and next classes we study the following things:
 - The *primary organization of databases* in storage and
 - How to physically store records in a file
 - The *techniques for accessing* them efficiently using various algorithms.
 - How to access the records more efficiently
 - Some algorithms require auxiliary data structures, called *indexes*.

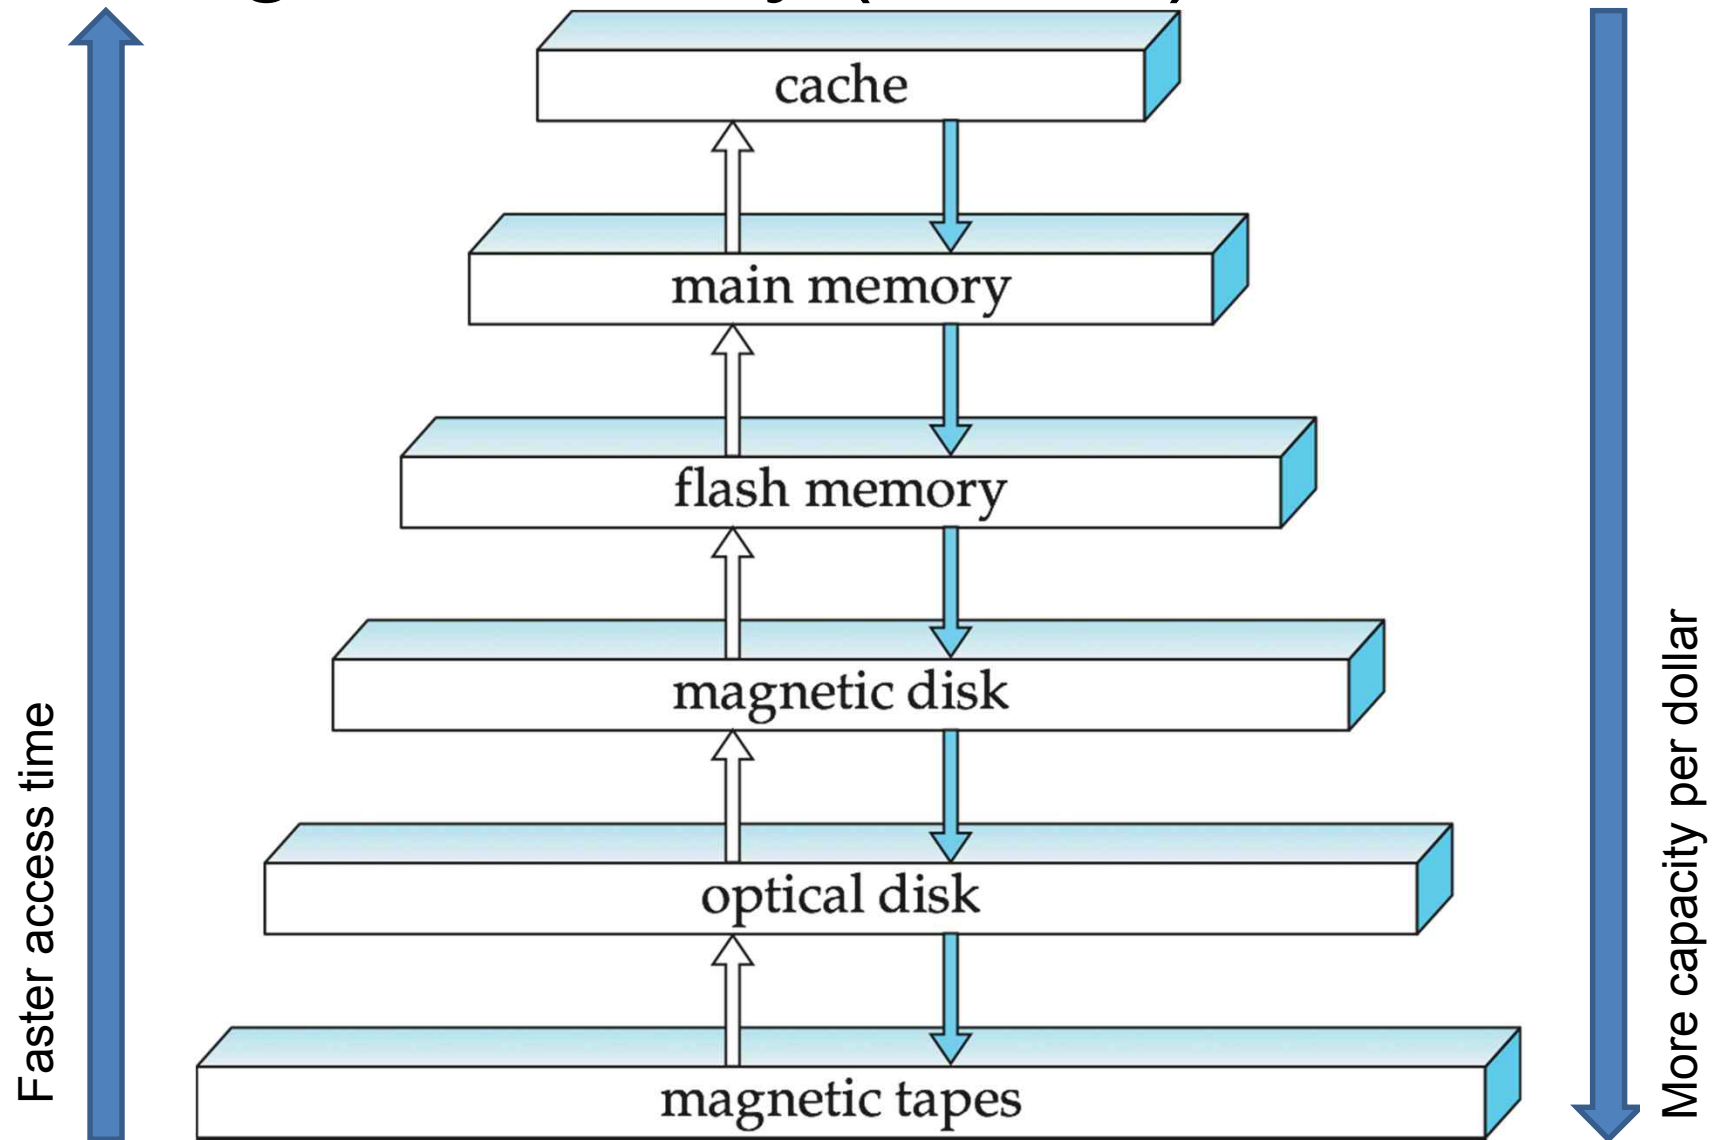
Introduction -

Storage Hierarchy

- The collection of data organizing database must be stored physically on some computer storage medium.
 - The DBMS (e.g., Oracle) software can then retrieve, update, and process this data as needed.
- Computer storage medium form a *storage hierarchy*:
 - **Primary storage**: operated by central processing unit (CPU)
 - CPU cache memory, RAM (random access memory),
 - **Secondary storage**: non-volatile, cheaper than primary
 - Hard disk drives (HDD), (USB) flash memory, solid-state drives (SSDs)
 - **Tertiary storage**: cheapest, but slowest
 - Removable media: tapes, optical disks (CD-ROMs, DVDs, ...)

Introduction -

Storage Hierarchy (Cont'd)



Introduction -

Storage Types and Characteristics

| Type | Capacity* | Access Time | Max Bandwidth | Commodity Prices (2014)** |
|-------------------------|------------------|-------------|---------------------|------------------------------|
| (Cache memory | 12MB | 0.5-2.5ns | 45GB/s | \$500-\$1000 per GB) |
| Main Memory- RAM | 4GB-1TB | 30ns | 35GB/sec | \$100-\$20K |
| Flash Memory- SSD | 64 GB-1TB | 50μs | 750MB/sec | \$50-\$600 |
| Flash Memory- USB stick | 4GB-512GB | 100μs | 50MB/sec | \$2-\$200 |
| Magnetic Disk | 400 GB-8TB | 10ms | 200MB/sec | \$70-\$500 |
| Optical Storage | 50GB-100GB | 180ms | 72MB/sec | \$100 |
| Magnetic Tape | 2.5TB-8.5TB | 10s-80s | 40-250MB/sec | \$2.5K-\$30K |
| Tape jukebox | 25TB-2,100,000TB | 10s-80s | 250MB/sec-1.2PB/sec | \$3K-\$1M+ |

*Capacities are based on commercially available popular units in 2014.

**Costs are based on commodity online marketplaces.



Magnetic tape



*IBM System Storage
TS3500 Tape Library,*

Introduction -

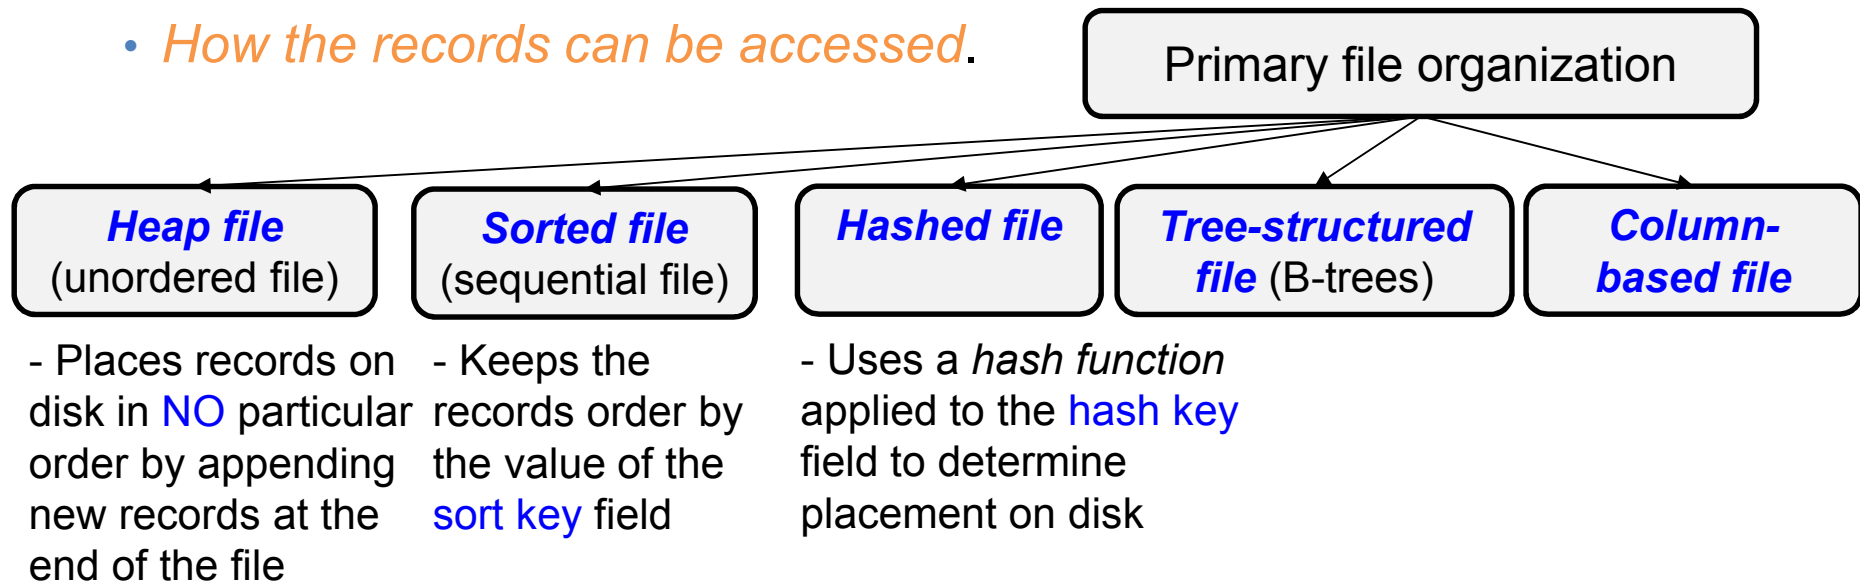
Storage Organization of Databases

- *Persistent data* (nonvolatile)
 - Most databases typically store large amounts of data persisting over long periods of time.
 - Stored on secondary storage on magnetic and/or SSD disks. Why?
 - 1) In general, DBs are **too large** to fit in main memory in its entirety.
 - 2) **Nonvolatile**; don't disappear after power off
 - 3) The cost of storage per unit of data: 10x **cheaper** than primary storage.
 - Portion of data is loaded from disk into RAM for processing and rewritten to the disk.
 - The data stored disk is organized as **files** of **records**.
 - Each record is a collection of data values representing facts about ER.
 - Records should be **well-placed** on disk for **efficient access**.
- C.f. *Transient data* (volatile): contrasts with persistent data.
 - Exist only during program execution; e.g., `malloc()`

Introduction -

Storage Organization of Databases (Cont'd)

- **Primary file organizations** determine:
 - How the file records are *physically placed* on the disk, *and hence*
 - *How the records can be accessed.*



- **Secondary organization** (or *auxiliary access structure*)
 - Allows efficient access to file records based on **alternate fields** different from those that have been used for the primary file organization.
 - Most of these exist **indexes**.

SECONDARY STORAGE DEVICES


Chapter 16.2

Disk Storage: Nonvolatile, rotating magnetic storage

- A disk: a random access addressable device
 - Transfer of data from/into RAM in unit of disk blocks
 - A **block** (or **sector**) **address** consists of
 - A cylinder number, a track number (within the platter), a block number (within the track) (or, a sector number (within the block))
 - The block address is supplied to the disk I/O controller.
 - In many modern disk drives, a single number called **LBA** (logical block address) is mapped automatically to the right block by the controller.
- *Disk **read***
 - The desired disk block with an LBA is copied into the disk drive **buffer**.
- *Disk **write***
 - The contents of the buffer is copied into the disk block with an LBA.

Disk Storage: Internal Architecture (Cont'd)

- *Disk* is a stack of magnetic *platters*.
 - The surface of each platter is organized into many *tracks*.
 - Tracks divided into *sectors*.
 - *Sector*: the basic unit of storage, consisting of
 - *Sector ID*, *data* (typically, 512 bytes), *Error-correction code*, *sync. fields*, or *gaps* (the sector # of the next sector)
 - *Block* (or *page*): a group of sectors (in unit of read/write), typically, 4K bytes
 - The disk heads move to the desired tracks (in the same *cylinder*) while the *spindle* rotates to locate desired sectors.
- * Most disk controllers with a *built-in cache* for better performance (of *what?*).




Disk Access Time

- Three major components
 - *Seek time*: taken to move the disk head to the desired track
 - *Rotational delay*: taken to wait for the desired sector to rotate until it comes under the (read/write) head
 - Usually assumed to *be half* (why is it so??) of the full rotation time
 - *Transfer time*: taken to transfer a disk block of bits
 - Function of the sector size, the rotation speed, and the recording density
 - 100 ~ 200 MB / sec in 2012
- Besides, there could be “*more latencies*” from *queuing delays* when other accesses exists and from *disk controller overhead*.
- Example: 512B sector, 7,200rpm, 4ms average seek time, 100MB/s transfer rate, 0.2ms controller overhead, idle disk
 - Q: What would be the expected “read” (or “write”) time of a sector on this disk (즉, 한 섹터 읽(쓰)는데 걸리는 평균 시간)?

Techniques for Efficient Data Access

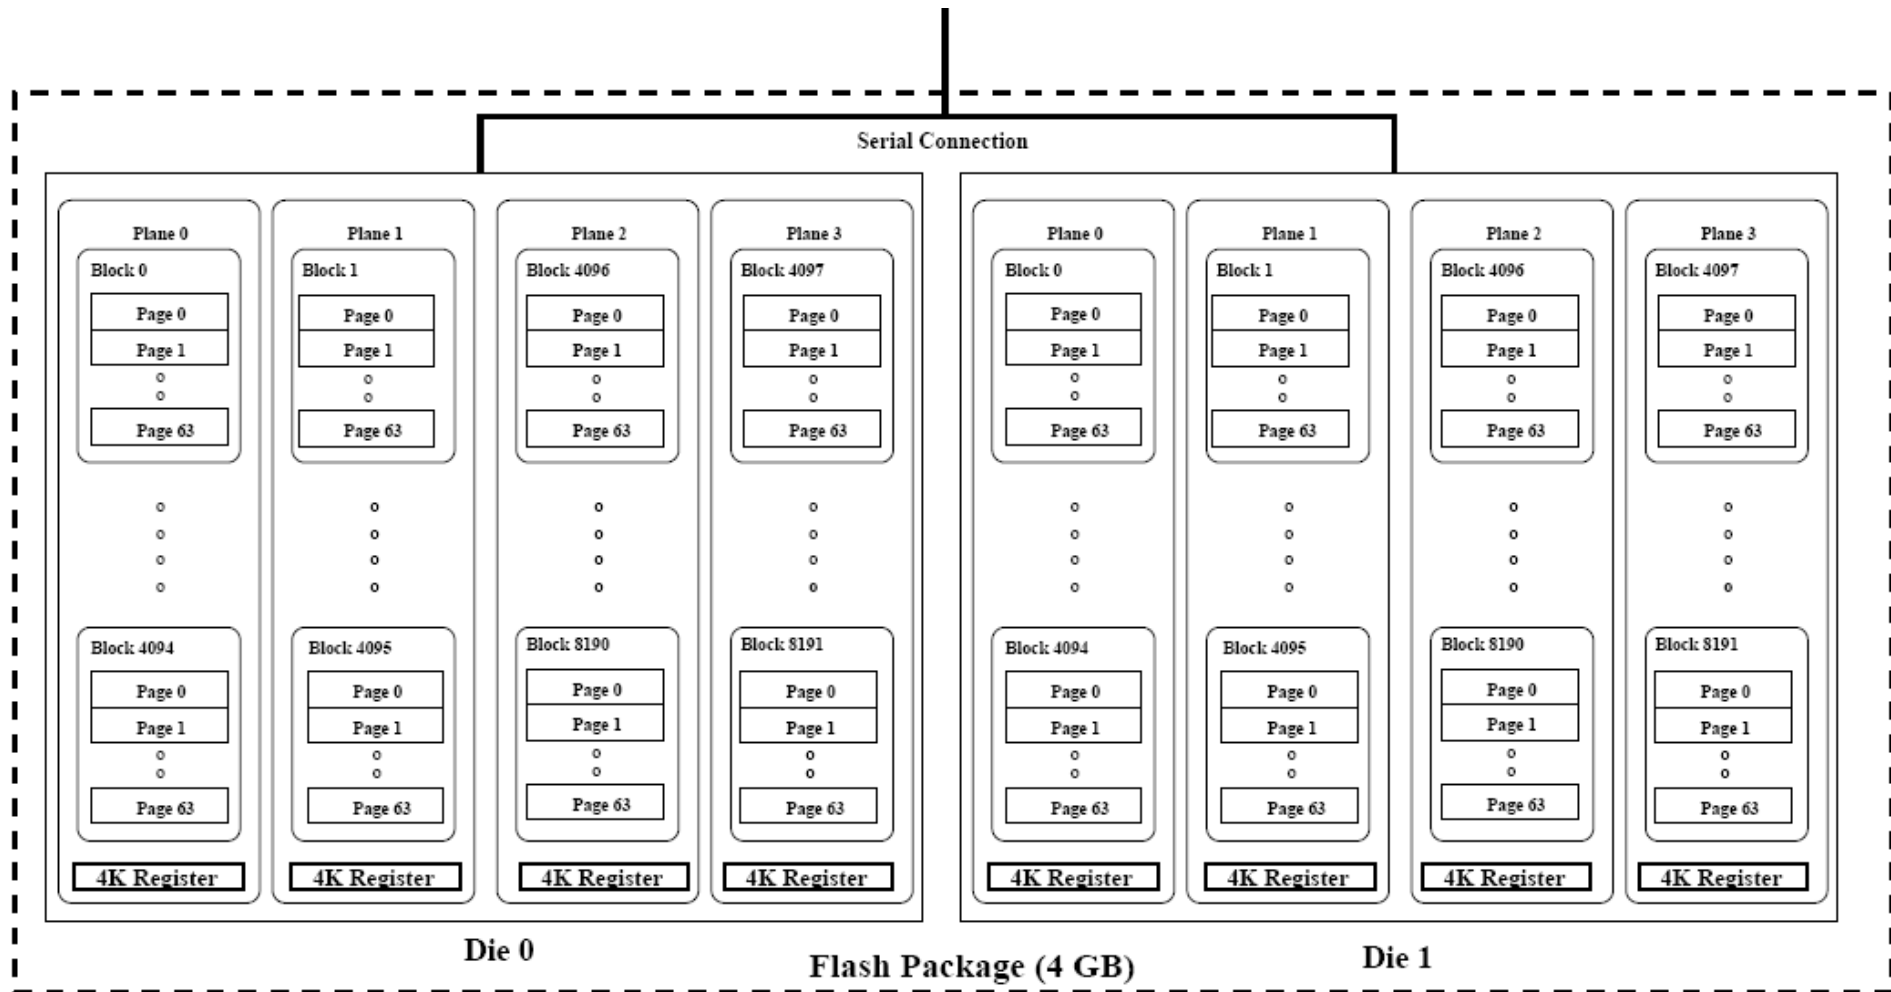
- Buffering of data (in main memory)
 - New data can be held in a buffer while old data is in processing.
- Proper organization of data on disk
 - Keep related data on contiguous blocks; place data blocks close to head
- Reading data ahead of request: called *prefetching*
 - For a disk read, blocks from the rest of the track can also be read.
- Proper scheduling of I/O requests
 - Aims at minimizing total access time;
 - Arms moves only in one direction: called *the elevator algorithm*
- Uses of log disks to temporarily hold writes
 - All blocks to be written go to one log disk to eliminate seek time.
- Use of SSDs: no latency of mechanical parts but expensive

Flash Memory



- Type of *Electrically Erasable Programmable Read-Only Memory* (EEPROM)
 - **Non-volatile** semiconductor storage with **NO** moving mechanical part
 - 100× – 1000× faster than disk: *msec (ms)* vs. *microsec (us)*
 - *Exceedingly fast* read speed, *smaller*, *less power*, *more robust*
 - But more \$/GB (between disk and DRAM)
- Two types over a flash cell
 - NOR-based flash: bit cell acts like a NOR gate
 - Providing (byte-addressable) random read/write access
 - More expensive, taking longer to erase and write new data
 - NAND-based flash: bit cell acts like a NAND gate
 - Denser (bits/area) thus more storage, but page(block)-addressable access
 - Cheaper per GB
- Disadvantages
 - *No in-place update*: asynchronous latency between read and write
 - *Limited lifetime*: cannot be written on one cell forever => *wear-leveling*

Flash Memory: SSD (Samsung 4GB)




Die (4 planes) – Plane (4K blocks) – Block (64 pages) – Page (2K (currently, 4K or 8K))

BUFFERING OF BLOCKS

Chapter 16.3

- *Buffer* refers to a part of main memory available to receive disk blocks.

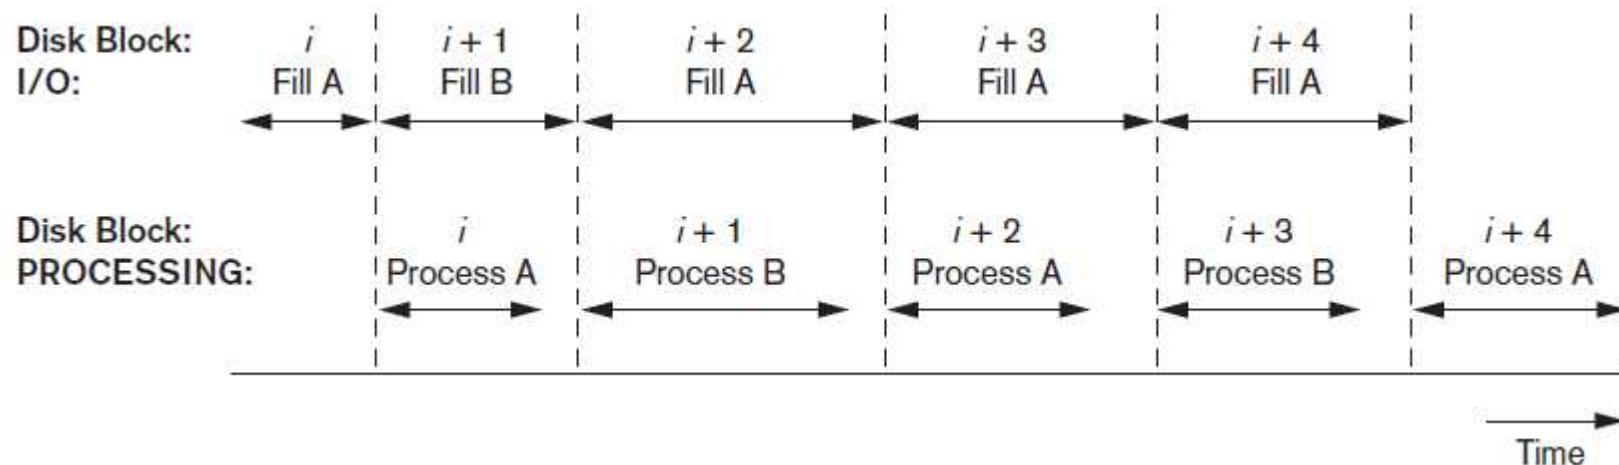
Interleaved Concurrency (on Single CPU) vs. Parallel Execution (on multi-CPU)



- Buffering is most useful when processes can run **concurrently in parallel**. Why? See next.
 - Because a separate disk I/O processor available or multiple CPU processors exist

Concept of *Double Buffering*


- Use of two buffers: **A** and **B**, for reading disk



- Illustrates how reading and processing can proceed when the time for processing is *smaller* the time for reading the next block (I/O time)
- **Double buffering:**
 - Once a block transfer is completed, CPU can start processing that block.
 - At the same time, the disk I/O controller (processor) can be reading and transferring the next block into a different buffer.
- Can be used to read continuous streams of blocks.

Buffer Management

- *Buffer manager* of a DBMS
 - Responds to data request
 - Decides (i) *which buffer to use* and (ii) *what pages to replace in the buffer* to accommodate new pages* (or, disk blocks*)
 - Views the available main memory storage as a **buffer pool**.
 - Keeps two types of information about each page:
 - 1) **Pin-count**: # of times that page has been requested (or # of concurrent users on that page); initially, 0
 - **Pinning**: incrementing the pin-count
 - If it falls to 0, then **unpinned**. Otherwise, its associated page cannot be evicted.
 - 2) **Dirty bit**: Initially, 0, but set to 1 whenever that page is updated by program



* Page and block are interchangeably used.

When a Certain Page is Requested...

- The buffer manager checks if the requested page is already in a buffer in the buffer pool.
 - If the page exists, then the manager increments its pin-count and releases the page.
 - If NOT, then the manager does the following things:
 - a)* It chooses a page for replacement, using the replacement policy (to be discussed shortly), and increments its pin-count.
 - b)* If the dirty bit of the replacement page is **ON**, the manager writes that page to disk (by replacing its old copy on disk). (If the bit is OFF, then no need to write back to the disk. Why?)
 - c)* It reads the requested page into the space just freed up.
 - d)* The main memory address of the new page is passed to the requestor.
- If there is no unpinned page and the requested page is not available in the buffer pool, then the manager must wait (until a page gets available).

Buffer Replacement Strategies

- Popular buffer replacement strategies
 - *Least recently used (LRU)*
 - Throw out that page that has not been used for the longest time.
 - *Clock policy*: a round-robin variant of LRU
 - Finds a buffer with a flag with a value of 0 in round-robin fashion
 - Assume each buffer can have a value of 0 (*unused*) or 1 (*used*).
 - First-in-first-out (FIFO)
 - Will replace the page that has been occupied the longest
 - [Caution] A root block of index may be thrown out, but

PLACING FILE RECORDS ON DISK

Chapter 16.4

Records and Record Type

- **Record**: collection of “related” data values or items
 - Values correspond to *record field* (or tuple attribute)

EMPLOYEE

| Fname | Minit | Lname | <u>Ssn</u> | Bdate | Address | Sex | Salary | Super_ssn | Dno |
|-------|-------|-------|------------|------------|--------------------------|-----|--------|-----------|-----|
| John | B | Smith | 123456789 | 1965-01-09 | 731 Fondren, Houston, TX | M | 30000 | 333445555 | 5 |

- **Record type (format)**
 - A collection of (i) field names and (ii) their corresponding data types


Ex) An EMPLOYEE record type in C-style notation

- Data types: numeric, string, boolean, date/time ...
- Binary large objects (BLOBs): unstructured objects (images, videos or audio stream)
 - A BLOB data item: stored separately from its record in a pool of disk blocks; a pointer to the BLOB kept in the record.
- Character large objects (CLOBs): for storing free text

```
struct Employee{
    char fname[10];
    char minit[1];
    char lname[10];
    char ssn[9];
    char bdate[10];
    char address[50];
    char sex[1];
    int salary;
    char superssn[9];
    short int dno;
};
```

Files, Fixed-Length Records, and Variable-Length Records

- File: a *sequence* of records.
- Fixed-length records
 - Every record in a file has exactly the same size (in bytes).
- Variable-length records
 - Different records in the file have different sizes.
- Why variable-length records? Four reasons.
 - 1) One or more fields have variable-length: e.g., VARCHAR
 - 2) One or more fields are repeating. e.g.
 - 3) One or more fields are optional.
 - 4) File contains records of different types.



| lastName | homePhone | workPhone | cellPhone | fax |
|----------|--------------|-----------|--------------|-----|
| Barnes | 562-874-1234 | | 310-999-3628 | |


Repeating!

TRANSCRIPT


| Student_name | Student_transcript | | | | |
|--------------|--------------------|-------|----------|------|------------|
| | Course_number | Grade | Semester | Year | Section_id |
| Smith | CS1310 | C | Fall | 08 | 119 |
| | MATH2410 | B | Fall | 08 | 112 |

3 Options for Formatting Records of a File of Variable-length Records: EMPLOYEE

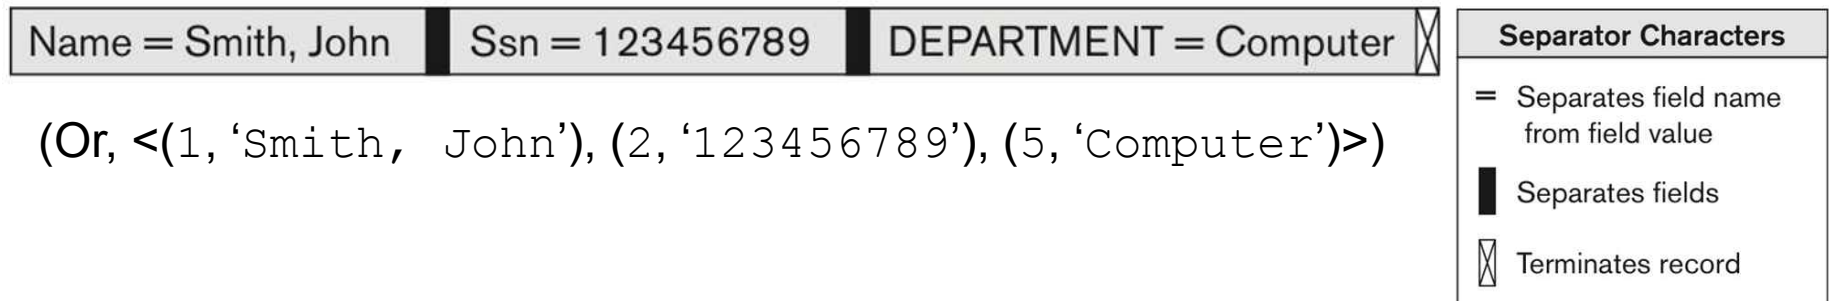
1) A fixed-length record with 6 fields and size of 71 bytes: easy location



2) A record with 2 variable-length fields and 3 fixed-length fields: for variable-length



3) A variable-field record with 3 types of separator characters: with *optional fields*




Record Blocking and Spanned vs. Unspanned Records

- The records of a file **MUST** be allocated to disk blocks.
 - Why? A block is the *unit of data transfer* between disk and memory.
 - When $|\text{block}| \geq |\text{record}|$, each block will contain numerous records.

- Spanned records*


- Larger than a single block (page: 4K or 8K)
- Pointer at end of first block points to block containing rest of record



Record 4: a spanned record

- Unspanned records*

- Used when a record is smaller than a block
- not allowed to cross block boundaries



Record Blocking and Spanned vs. Unspanned Records (Cont'd)

- *Blocking factor*: avg. # of records per block for the file
 - Some unused space = $B - (bfr * R)$ bytes, where
 - B : block size, R : record size, and bfr : blocking factor for the file
 - bfr can be used to calculate # of blocks, say b , needed for a file of r records.
 - $b = \text{ceiling}(r / bfr)$ blocks, where
 - r : the number of records for the file
- Q: How many blocks are needed to store a file of 1000 records?
 - Record size: 16 Bytes
 - Block size: 4 KBytes
 - Unused space: 0 Bytes (for convenience)

Allocating File Blocks on Disk

- Several standard techniques
 - **Contiguous allocation**: the file blocks are allocated to consecutive disk blocks
 - The whole file can be read fast by double buffering while being hard to make it expanded.
 - **Linked allocation**: each file block contains a pointer to the next file block.
 - Easy to expand but make it slow to read the whole file.
 - **Cluster allocation**: a combination of the two
 - Allocates clusters, or called *file segments* or *extents*, which are linked.
 - **Indexed allocation**: one or more index blocks contain pointers to the actual file blocks. (e.g., a hash file)

File Headers

- A **file header** (descriptor) contains information about a file needed by the system programs accessing the file records.
 - Which information?
 - To determine the disk addresses of the file blocks
 - To record format descriptions: e.g., field-lengths, the order of fields within a record for fixed-length unspanned records and field type codes, separator characters, and record type codes for variable-length records

Database file

| | | | | |
|-----------------|-----------|---------------|-----------|-----|
| <i>header</i> | | | | |
| <i>record 1</i> | 123456789 | John Smith | New York | 5 |
| <i>record 2</i> | 234567891 | Chris Young | San Diego | 4 |
| | .. | ... | ... | ... |
| <i>record n</i> | 987654321 | Christian Lee | Tucson | 1 |

OPERATIONS ON FILES

Chapter 16.5

Operations on Files

- *Retrieval* operations
 - Do not change any data in the file
 - Only locate certain records so that their field values can be examined and processed
- *Update* operations
 - Change the file by insertion, or deletion of records or by modification of field values.
- In either case, we should select one or more records
 - Based on a selection (or filtering) condition
 - Specifying criteria that the desired record(s) must satisfy

File Organization (vs. Access Method)

- *File organization*
 - Refers to the organization of the data of a file into records, blocks, and access structures.
 - Includes the way records and blocks are placed on the disk and interlinked.
- How to organize records in files? (To be discussed in detail)
 - *Heap*: a record can be placed anywhere in the file where there is space
 - *Sequential*: stores records in sequential order, based on the value of the search key of each record
 - *Hashing*: uses a hash function computed on some attribute of each record; the result specifies in which block of the file the record should be placed
 - *Multitable clustering file organization*: records of each relation may be stored in a separate file.
 - A records of several different relations can be stored in the same file.
 - Motivation: store related records on the same block to minimize I/O

(File Organization vs.) **Access Method**

(Cont'd)

- *Access method*
 - Provides a group of operations that can be applied to a file:
 - Open, Find, FindNext, Read, Delete, Modify, Insert, Close, Scan.
 - Possible to apply several access methods to a file organized using a certain organization.
 - Some access methods can be applied only to files organized in certain ways.
 - Ex) An indexed access method can't be applied to a file without an index.

FILES OF UNORDERED RECORDS (**HEAP FILES**)


Chapter 16.6

Head (or Pile) File

- Simplest and most basic type of organization
- Records are placed in the file in the order in which they are inserted.
 - New records are inserted at the end of the file.
- “Insertion” of a new record is **very efficient**.
 - The last disk block of the file is copied into a buffer.
 - The new record is added.
 - The block is then rewritten back to disk.
 - The address of the last file block is kept in the file buffer.
- But “searching” a record is **inefficient** due to a **linear search**.
 - On average, $(b / 2)$, where b : # of blocks in a file.
 - In the worst, b file blocks will be visited.

Head (or Pile) File (Cont'd)

- One way to “delete” a record, a program must
 - First find its block,
 - Copy the block into a buffer,
 - Delete the record from the buffer, and finally,
 - Rewrite the block back to the disk.
 - Another way: to use a **deletion marker**
 - An extra bit or byte is stored with each record.
 - A record is deleted by setting the marker to a certain value.
 - A different value for the marker indicates a valid record.
 - Search considers only valid records in a block.
- ⇒ Both approaches required periodic **reorganization** of the file to reclaim the unused space of deleted records; packing of existing undeleted records
- Third way: to use the space of deleted records for insertion



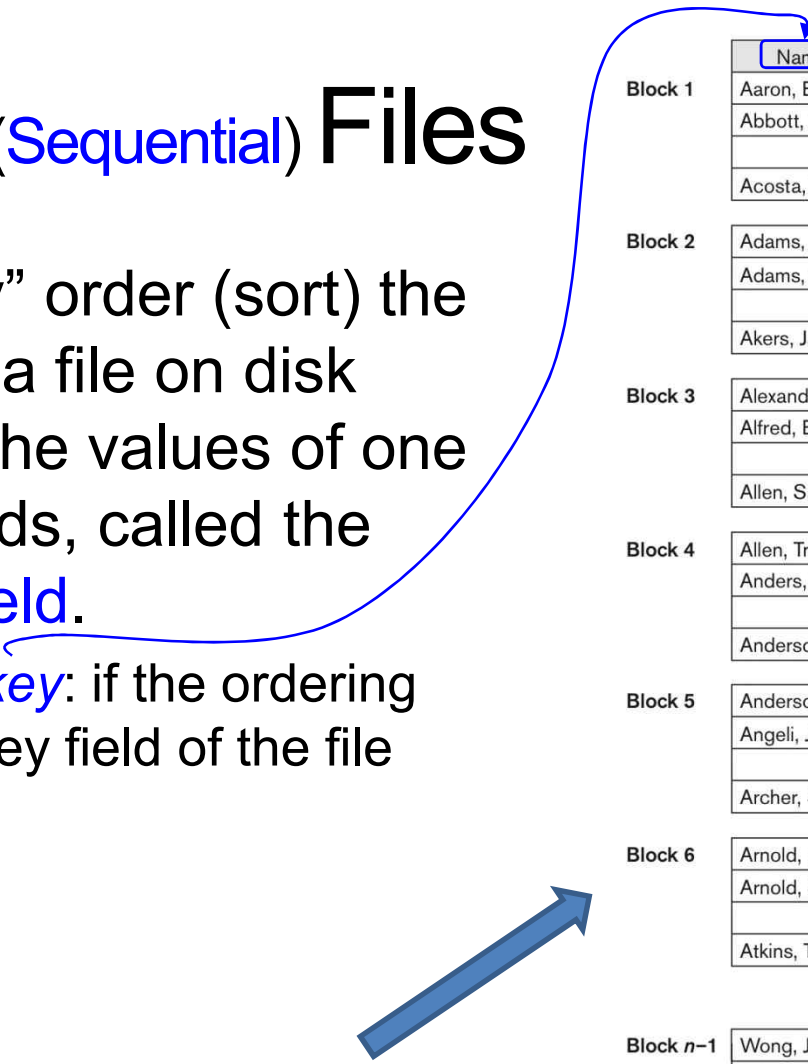
FILE OF ORDERED RECORDS (**SORTED FILES**)

Chapter 16.7

Sorted (Sequential) Files

- “Physically” order (sort) the records of a file on disk based on the values of one of their fields, called the **ordering field**.
 - Ordering key**: if the ordering field is a key field of the file

Illustrates an ordered file with Name as the ordering key field (assuming names are distinct.)



| | Name | Ssn | Birth_date | Job | Salary | Sex |
|-------------|-----------------|-----|------------|-----|--------|-----|
| Block 1 | Aaron, Ed | | | | | |
| | Abbott, Diane | | | | | |
| | ⋮ | | | | | |
| | Acosta, Marc | | | | | |
| Block 2 | Adams, John | | | | | |
| | Adams, Robin | | | | | |
| | ⋮ | | | | | |
| | Akers, Jan | | | | | |
| Block 3 | Alexander, Ed | | | | | |
| | Alfred, Bob | | | | | |
| | ⋮ | | | | | |
| | Allen, Sam | | | | | |
| Block 4 | Allen, Troy | | | | | |
| | Anders, Keith | | | | | |
| | ⋮ | | | | | |
| | Anderson, Rob | | | | | |
| Block 5 | Anderson, Zach | | | | | |
| | Angeli, Joe | | | | | |
| | ⋮ | | | | | |
| | Archer, Sue | | | | | |
| Block 6 | Arnold, Mack | | | | | |
| | Arnold, Steven | | | | | |
| | ⋮ | | | | | |
| | Atkins, Timothy | | | | | |
| ⋮ | | | | | | |
| | | | | | | |
| | | | | | | |
| | | | | | | |
| Block $n-1$ | Wong, James | | | | | |
| | Wood, Donald | | | | | |
| | ⋮ | | | | | |
| | Woods, Manny | | | | | |
| Block n | Wright, Pam | | | | | |
| | Wyatt, Charles | | | | | |
| | ⋮ | | | | | |
| | Zimmer, Byron | | | | | |

Sorted (Sequential) Files (Cont'd)

- Advantages (over heap files)
 - Reading records by ordering key is extremely efficient.
 - **NO sorting required**; strong over a search/range condition on the key
 - Finding the next record requires **no additional block access** unless the record is the last one in the block.
 - Allows the **binary search technique** to be used for performing a search over the value of an ordering key, resulting in “faster access”.

```
l = 1; u = b; //b: # of blocks
while (u >= l) {
    i = (l+u)/2;
    read block i of the file
    into the buffer.
    if K < (ordering key field
           value of the first
           record in block i)
        then u = i-1;
    else if K > (ordering key
                field value of
                the last record
                in block i)
        then l = i+1;
    else if K == (ordering key
                 field value)
        then break; // found
    else break; // not found
}
```

Sorted (Sequential) Files (Cont'd)

- Deletion: utilize **pointer chains** or deletion marker

Ordering key

- Insertion: locate the position where the record is to be inserted

- If there is free space, insert there
- If no free space, insert the record in an **overflow** (or transaction) **block**
- In either case, pointer chain must be updated.

| | | | | |
|-------|------------|------------|-------|--|
| 10101 | Srinivasan | Comp. Sci. | 65000 | |
| 12121 | Wu | Finance | 90000 | |
| 15151 | Mozart | Music | 40000 | |
| 22222 | Einstein | Physics | 95000 | |
| 32343 | El Said | History | 60000 | |
| 33456 | Gold | Physics | 87000 | |
| 45565 | Katz | Comp. Sci. | 75000 | |
| 58583 | Califieri | History | 62000 | |
| 76543 | Singh | Finance | 80000 | |
| 76766 | Crick | Biology | 72000 | |
| 83821 | Brandt | Comp. Sci. | 92000 | |
| 98345 | Kim | Elec. Eng. | 80000 | |
| | | | | |
| 32222 | Verdi | Music | 48000 | |

- Need to reorganize the file from time to time to restore sequential order


Sorted (Sequential) Files (Cont'd)

- Average access times for a file of b blocks under basic file organizations

| Type of Organization | Access/Search Method | Average Blocks to Access a Specific Record |
|----------------------|---------------------------------|--|
| Heap (unordered) | Sequential scan (linear search) | $b/2$ |
| Ordered | Sequential scan | $b/2$ |
| Ordered | Binary search | $\log_2 b$ |

APPENDIX


Another Disk Internal



(a) A single-sided disk with read/write hardware.

(b) A disk pack with read/write hardware.

Different Sector Organizations on Disk



(a) Sectors subtending a fixed angle.

(b) Sectors maintaining a uniform recording density.