



# QUERY OPTIMIZATION

---

## Chapter 19

Prof. Young-Kyoon Suh

# Introduction

- Query optimization(질의 최적화)
  - Conducted by a *query optimizer* in a DBMS
  - Goal
    - Select the *best available strategy* for query evaluation
    - Arrive *within a reasonable amount of time* at the *most efficient and the cost-effective plan* using the available information about the schema and the content of relations involved
- Most RDBMSs use a *tree* as the internal representation of a query.

# QUERY TREES AND HEURISTICS FOR QUERY OPTIMIZATION

---

## Chapter 19.1

# Heuristics-based Optimization Techniques

- Step 1: the scanner and parser of a query generates a data structure representing an *initial query tree presentation*
- Step 2: representation is *optimized* according to **heuristic rules**.
  - Leads to an *optimized query representation*, corresponding to the query execution strategy
  - One of the main heuristic rules: to apply SELECT and PROJECT *before* JOIN or other binary operations.
    - Why? To reduce size of files to be joined.
- Step 3: a *query execution plan* is generated.
  - The plan executes groups of operations based on the access paths available on the files involved in the query

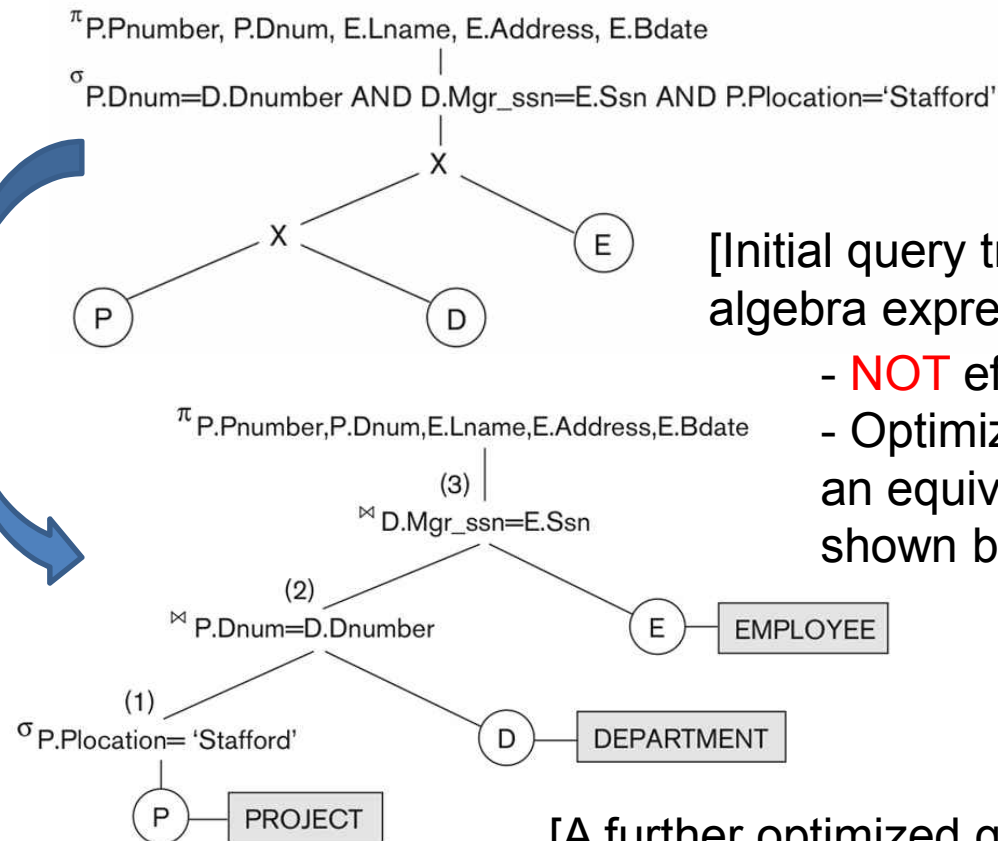
# Query (Evaluation) Tree

- A tree data structure corresponding to an extended relational algebra expression
  - *Leaf nodes*: *input relations* of the query
  - *Internal nodes*: *relational algebra operations* (or, *query operators*)
- An execution of the query tree consists of:
  - Executing a *query operator node* whenever its operands are available
  - Replacing the node by *the resulting relation* produced by the operator
- The order of the execution:
  - *Starts at the leaf nodes*, or the input relation
  - *Ends at the root node*, or the final operation of the query

# Example of a Query Tree

- Q1: `SELECT P.Pnumber, P.Dnum, E.Lname, E.Address, E.Bdate`  
`FROM PROJECT P, DEPARTMENT D, EMPLOYEE E`  
`WHERE P.Dnum=D.Dnumber AND D.Mgr_ssn=E.Ssn AND P.Plocation='Stafford'`

Applying  
the main  
heuristic  
rule: called  
**early  
selection**



[Initial query tree based on relational algebra expressions (RAEs)]

- **NOT** efficient; never executed
- Optimizer will transform into an equivalent final query tree shown below.

[A further optimized query tree based on RAE]

# Query Transformation Example

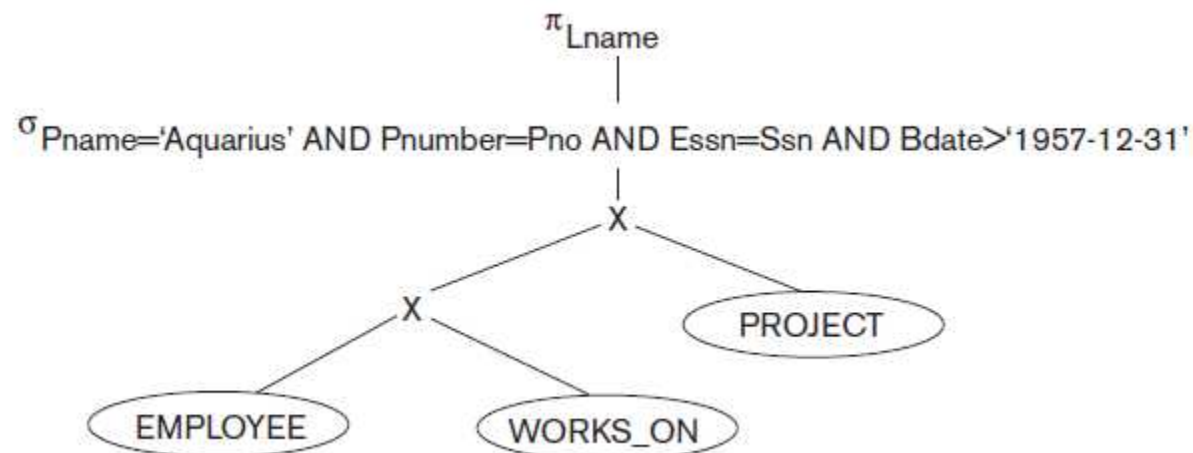
- Many different query trees can be *semantically equivalent*;
  - They represent *the same query* and *produce the same results*.

Q:   SELECT   E.Lname  
      FROM    EMPLOYEE E, WORKS\_ON W, PROJECT P  
      WHERE   P.Pname='Aquarius' AND P.Pnumber=W.Pno AND E.Essn=W.Ssn  
              AND E.Bdate > '1957-12-31';

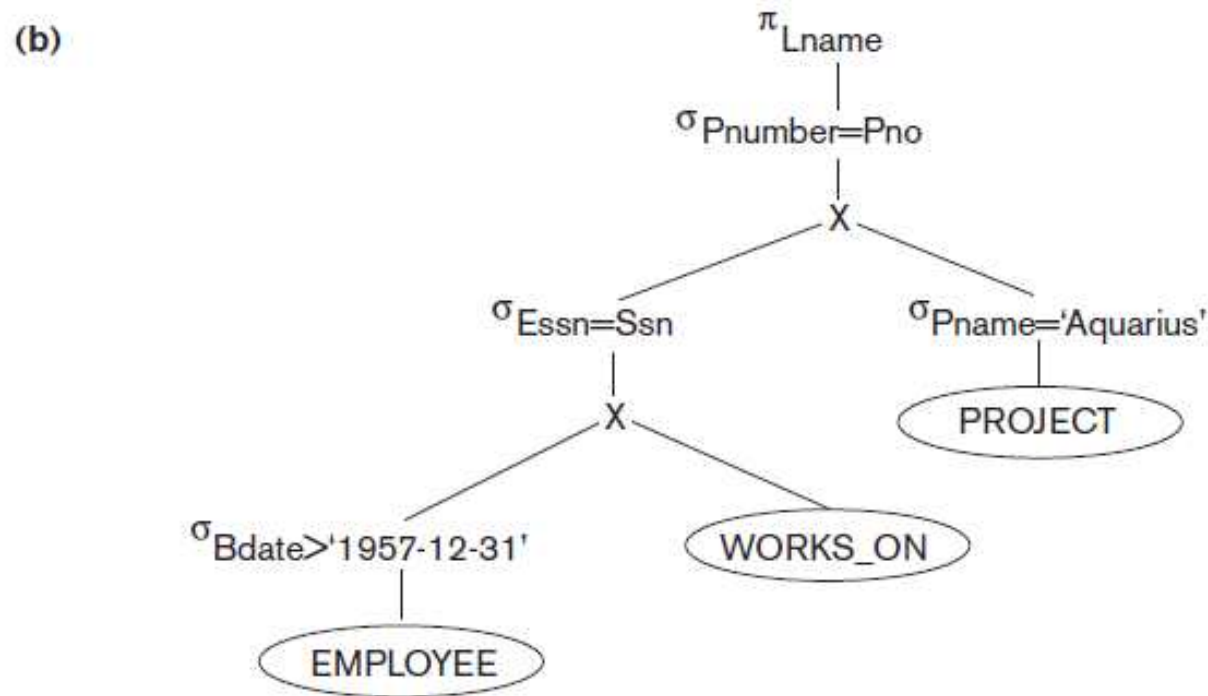


Converts to an initial query tree

(a)



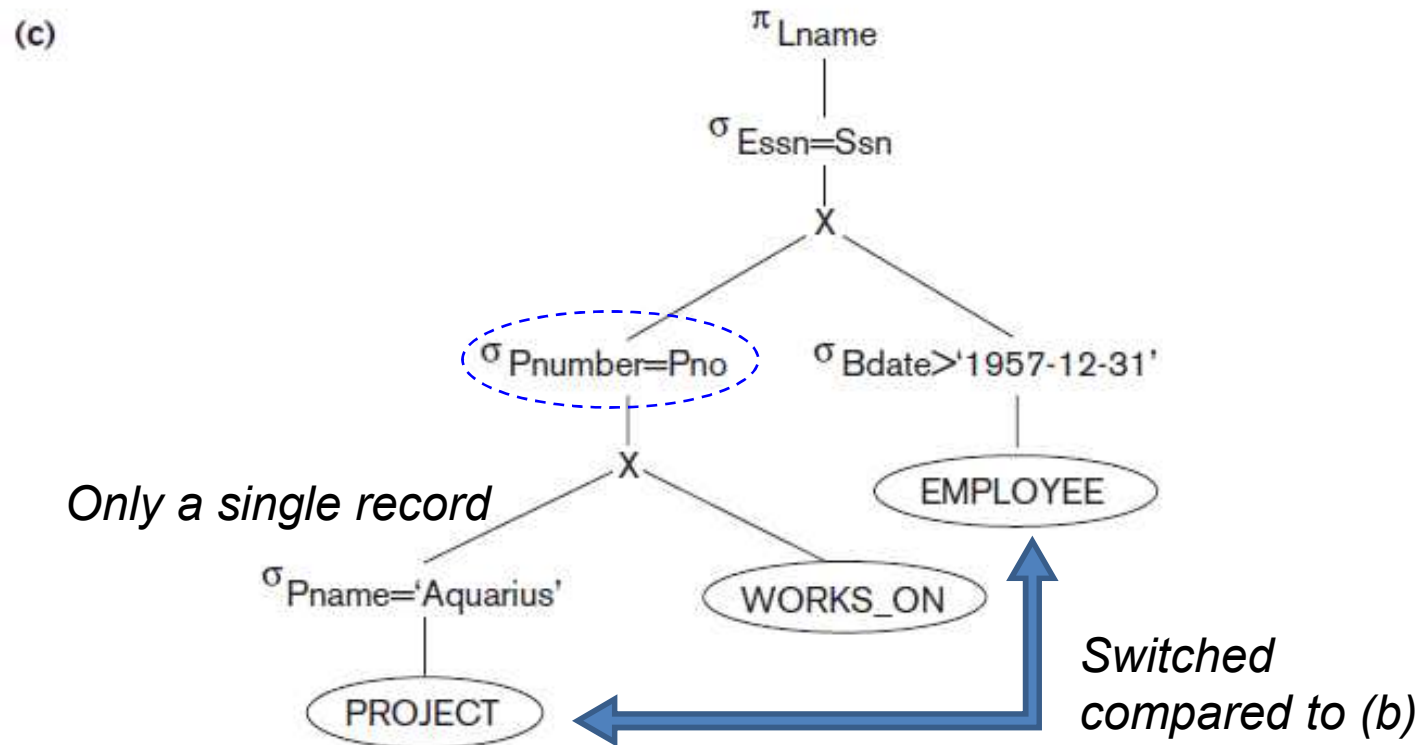
# Query Transformation Example (Cont'd)



[Moving SELECT operations down the query tree]

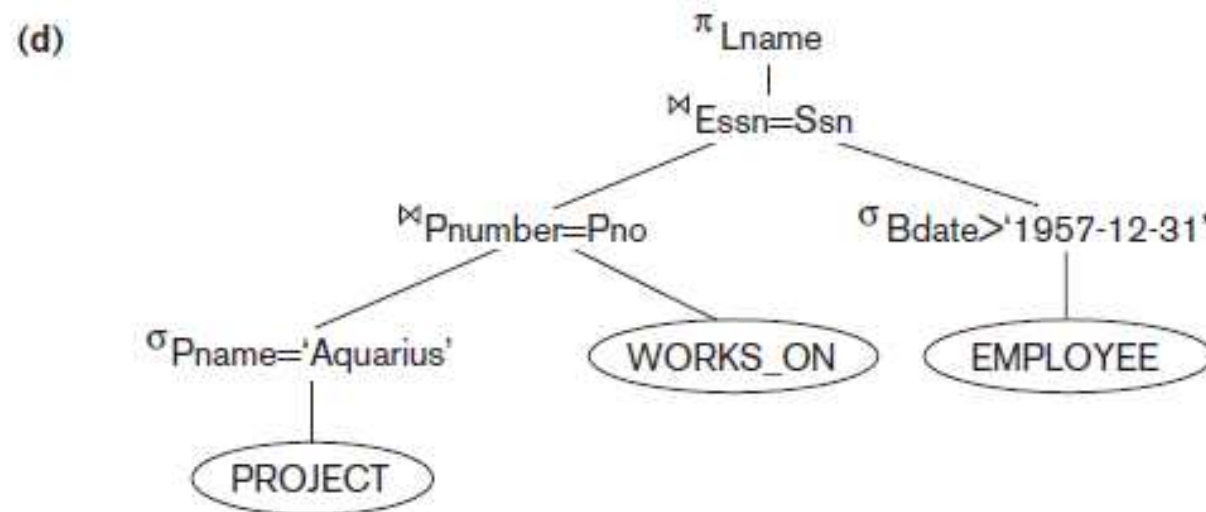


# Query Transformation Example (Cont'd)



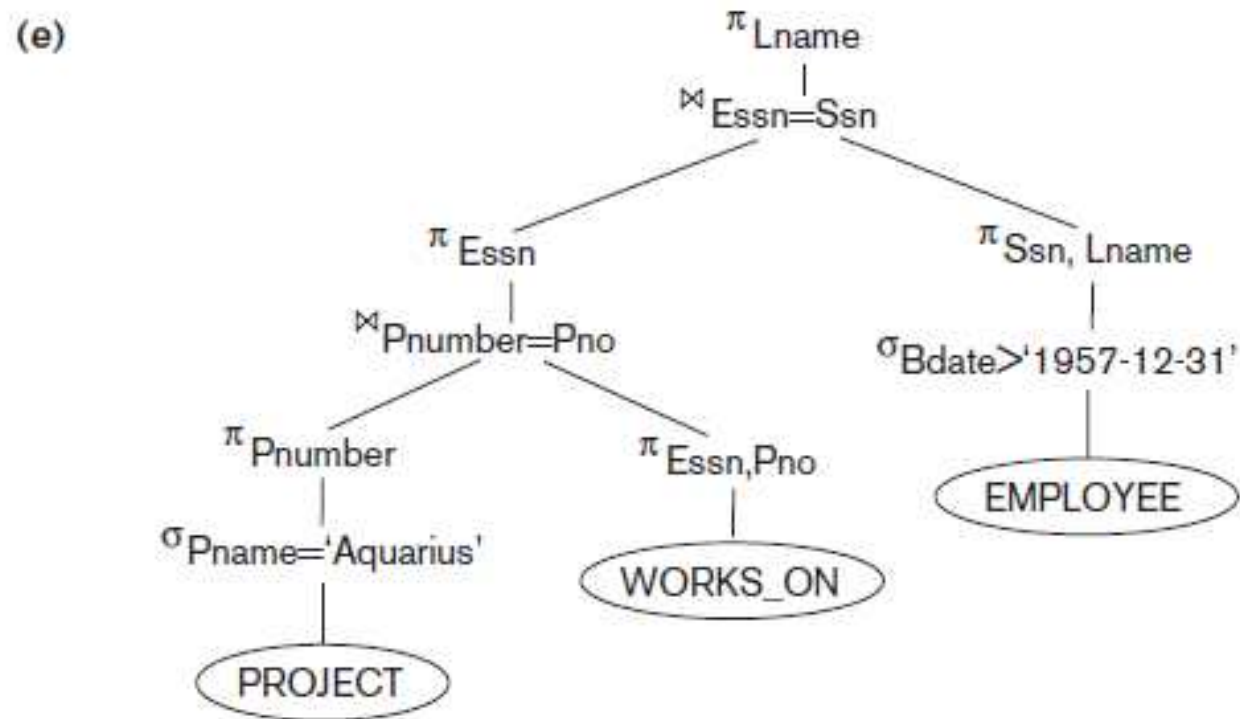
[Applying the *more restrictive* SELECT operation first]

# Query Transformation Example (Cont'd)



[Replacing CARTESIAN PRODUCT and SELECT with **JOIN** operations]

# Query Transformation Example (Cont'd)



[Moving PROJECT operations down the query tree]

# Summary of Heuristics for Algebraic Optimization

- Apply first operations that *reduce the size of intermediate results*
  - Perform SELECT and PROJECT operations **as early as possible** to reduce the number of tuples and attributes
  - The SELECT and JOIN operations that are most restrictive should be **executed before** other similar operations.
    - So that fewer intermediate results can be joined if any

# COST-BASED OPTIMIZATION

---

## Chapter 19.3

# *Cost-based Query Optimization* (QO)

- Applied to *compiled queries* rather than *interpreted queries*.
- Refers to the following approach:
  - An optimizer *enumerates* possible *execution plans* for a given SQL query.
  - It then *estimates and compares execution costs* for a query using different execution strategies and algorithms.
  - It then *chooses the execution plan with the lowest cost*.
- The scope of QO: a single query block
  - The optimizer consider various information: *various table and index access paths, join orders, join methods, group-by methods*, etc.
- Note: a query optimizer does *not* depend *solely* on heuristic rules or query transformations.

# Cost Components for Query Execution

- *Disk I/O cost*
    - Access cost to secondary storage
  - *Space (disk storage) cost*
    - Cost of storing on disk any intermediate results
  - *CPU (computation) cost*
    - Cost of performing in-memory operations on records within buffer
  - *Memory usage cost*
    - Cost pertaining to # buffers needed for query execution
- (*Communication cost*  
- *query shipping and receiving results from a remote site*)

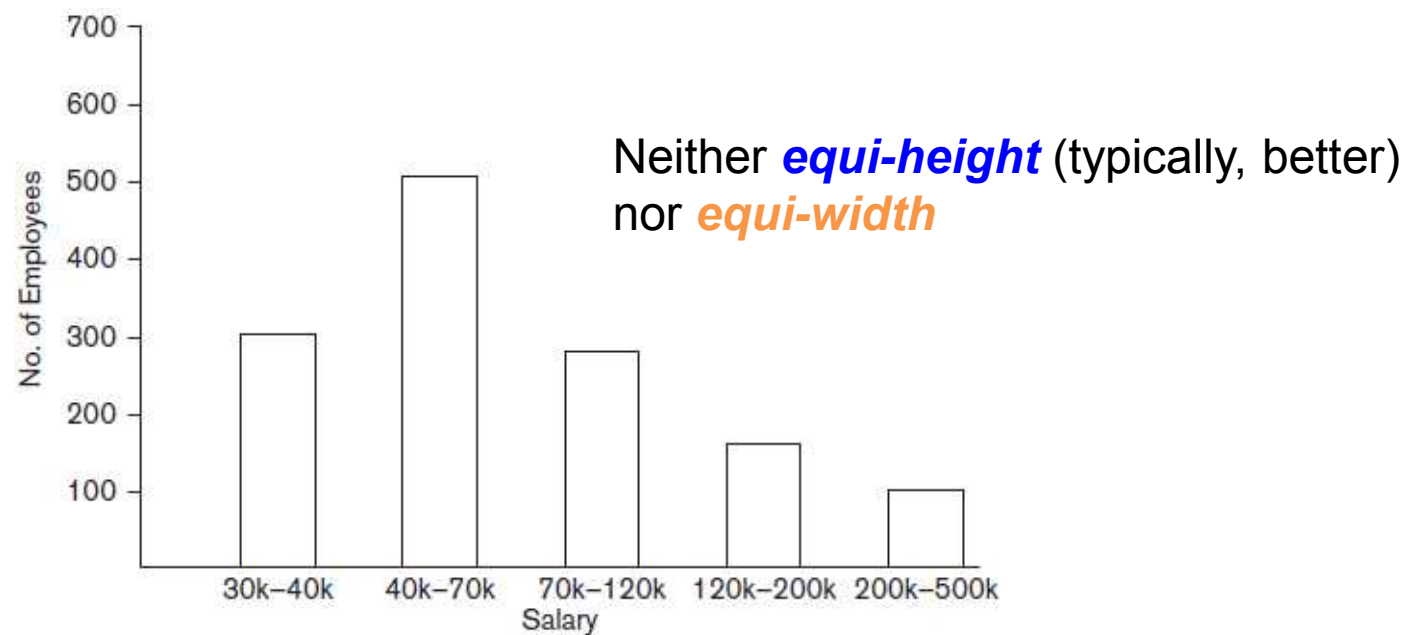
# Catalog Information Used in Cost Functions

- To estimate the costs of various plans, we leverage *cost functions*, and any information needed for the functions must be kept.
- The catalog information used by query optimizer
  - 1) File size
    - # records, the (average) record size, # blocks for a relation file, *bfr*
  - 2) File organizations
    - Primary organizations: unordered, ordered, hashed, tree-structured, etc.
      - Primary, clustering indexes with their indexing attributes
      - Secondary indexes with their indexing attributes , etc.
  - 3) Number of levels of each multi-level index
  - 4) Number of distinct values of an attribute
  - 5) Attribute selectivity: avg. # records satisfying equality condition



# Histograms

- Tables or data structures that record information about *data distribution*
- RDBMS stores histograms for most important attributes



[An Example of Histogram of salary in the relation EMPLOYEE]

# COST FUNCTIONS FOR SELECT OPERATION

---

Chapter 19.4

# Notations Used in Cost Formulas

$C_{Si}$ : Cost for method  $Si$  in block accesses

$r_X$ : Number of records (tuples) in a relation  $X$

$b_X$ : Number of blocks occupied by relation  $X$  (also referred to as  $b$ )

$bfr_X$ : Blocking factor (i.e., number of records per block) in relation  $X$

$sl_A$ : Selectivity of an attribute  $A$  for a given condition

$s_A$ : Selection cardinality of the attribute being selected ( $= sl_A * r$ )

$x_A$ : Number of levels of the index for attribute  $A$

$b_{11}A$ : Number of first-level blocks of the index on attribute  $A$

$NDV(A, X)$ : Number of distinct values of attribute  $A$  in relation  $X$

# Cost Functions for SELECT

- S1: *Linear search* (brute force approach)

- Search all file blocks to retrieve all records

$$C_{S1a} = b$$

- For equality condition on a key attribute, on average *one-half* the records are searched

$$C_{S1b} = \frac{b}{2}$$

- S2: *Binary search*

$$C_{S2} = \log_2 b + \left\lceil \frac{s}{bfr} \right\rceil - 1$$

- Reduces to  $\log_2 b$  if equality condition is on a key attribute
  - $s$ : selection cardinality

# Cost Functions for SELECT (Cont'd)

- S3a: Using a *primary index* to retrieve a single record

$$C_{S3a} = x + 1$$

- S3b: Using a *hash key* to retrieve a single record

$$C_{S3b} = 1$$

- S4: Using an *ordering index* to retrieve multiple records

$$C_{S4} = x + \frac{b}{2}$$

- S5: Using a *clustering index* to retrieve multiple records

$$C_{S5} = x + \left\lceil \frac{s}{bfr} \right\rceil$$

- S6: Using a *secondary(B+ tree) index* to retrieve multiple records

$$C_{S6a} = x + 1 + s \text{ (worst case), } C_{S6b} = x + \frac{b_{11}}{2} + \frac{r}{2} \text{ (for range queries)}$$

# Cost Functions for SELECT (Cont'd)

- **Dynamic programming** may be considered to find the optimal plan (with the least execution cost) without considering all possible execution plans.
  - Cost-based optimization approach
  - Sub-problems are solve only once.
  - Applicable when a problem may be broken down into subproblems that themselves have subproblems.
    - E.g. The cost of a root can be broken into the cost of each subtree of the tree, etc.

# COST FUNCTIONS FOR THE JOIN OPERATION

---

Chapter 19.5

# Cost Functions for JOIN

- Cost functions involve estimate of file size that results from the JOIN operation
  - Recall join may leave intermediate results during processing, leading to nontrivial I/O cost
- *Join selectivity* ( $js$ )
  - Ratio of the size of resulting file to size of CARTESIAN PRODUCT file
  - Simple formula for join selectivity on  $R \bowtie_{R.A=S.B} S$ 
    - $js = \frac{1}{m \max(NDV(A,R), NDV(B,S))}$ , when  $A$  or  $B$  is a key of either  $R$  or  $S$ , each.
- *Join cardinality* ( $jc$ )
  - The size of the resulting file after join:  $jc = js \times |R| \times |S|$



# Cost Functions for JOIN (Cont'd)

- J1: ***Nested-loop join***

- For three memory buffer blocks:

$$C_{J1} = b_R + (b_R * b_S) + ((js * |R| * |S|)/bfr_{RS})$$

- For  $n_B$  memory blocks:

$$C_{J1} = b_R + (\lceil b_R/(n_B - 2) \rceil * b_S) + ((js * |R| * |S|)/bfr_{RS})$$

- J2: ***Index-based nested-loop join***

- For a secondary index with selection cardinality  $S_B$  for join attribute  $S.B$

$$C_{J2a} = b_R + (|R| * (x_B + 1 + s_B)) + ((js * |R| * |S|)/bfr_{RS})$$

$((js * |R| * |S|)/bfr_{RS})$ : the cost of writing results to disk

# Cost Functions for JOIN (Cont'd)

- J3: **Sort-merge join**

- For files already sorted on the join attributes (without external sort)

$$C_{J3a} = b_R + b_S + ((js * |R| * |S|)/bfr_{RS})$$

- Cost of sorting must be added if sorting needed like external sort

- J4: **Partition-hash join**

$$C_{J4} = \boxed{3 * (b_R + b_S)} + ((js * |R| * |S|)/bfr_{RS})$$

$R$  and  $S$  are read twice for partitioning via hashing.  
Finally, partitioned  $R$  and  $S$  are written once.

# Example Table for Calculating Execution Cost

Cost Estimate Items	Number
$r_{E(MPLOYEE)}$	10,000 records
$b_{E(MPLOYEE)}$	2,000 blocks
$r_{D(DEPARTMENT)}$	125 records
$b_{D(DEPARTMENT)}$	13 blocks
$x_{Dnumber}$ (Number of levels of a primary index on the <code>Dnumber</code> attribute in <code>DEPARTMENT</code> )	1
$x_{Dno}$ (Number of levels of a secondary index on the <code>Dno</code> attribute in <code>EMPLOYEE</code> )	2
$s_{Dno}$ (Selectivity cardinality on <code>Dno</code> )	80 (records selected)
Join selectivity ( $js$ ) on $EMPLOYEE \bowtie_{Dno=Dnumber} DEPARTMENT$	$\frac{1}{\max(NDV(Dno, EMPLOYEE), NDV(Dnumber, DEPARTMENT))} = \frac{1}{125}$
Blocking factor of the resulting join ( $bfr_{ED}$ )	4 records

`DEPARTMENT.Dnumber`: a key attribute of `DEPARTMENT`

# Example of Join Optimization Based on Cost Functions

$\text{EMPLOYEE (E)} \bowtie_{\text{DNO=Dnumber}} \text{DEPARTMENT (D)}$

1) Nested-Loop Join (NLJ) with EMPLOYEE as outer loop:

- $$C_{J1} = b_E + (b_E * b_D) + ((js * r_E * r_D) / bfr_{ED})$$
$$= 2,000 + (2,000 * 13) + (((1/125) * 10,000 * 125) / 4) = 30,500$$

2) Nested-Loop Join (NLJ) with DEPARTMENT as outer loop:

- $$C_{J1} = b_D + (b_D * b_E) + ((js * r_E * r_D) / bfr_{ED})$$
$$= 13 + (13 * 2,000) + (((1/125) * 10,000 * 125) / 4) = 28,513$$

# Example of Join Optimization Based on Cost Functions (Cont'd)

$\text{EMPLOYEE (E)} \bowtie_{\text{DNO=Dnumber}} \text{DEPARTMENT (D)}$

## 3) Indexed-based NLJ with EMPLOYEE as outer loop:

- Using the primary index on DEPARTMENT in inner loop
- $C_{J2c} = b_E + (r_E * (x_{\text{Dnumber}} + 1)) + ((js * r_E * r_D) / bfr_{ED})$   
 $= 2,000 + (10,000 * (1 + 1)) + ((1/125) * 10,000 * 125) / 4 = 24,500$

## 4) Indexed-based NLJ with DEPARTMENT as outer loop:

- Using the secondary index on EMPLOYEE in inner loop
- $C_{J2a} = b_D + (r_D * (x_{\text{Dno}} + s_{\text{Dno}})) + ((js * r_E * r_D) / bfr_{ED})$   
 $= 13 + (125 * (2 + 80)) + ((1/125) * 10,000 * 125) / 4 = 24,500$

# Example of Join Optimization Based on Cost Functions (Cont'd)


EMPLOYEE (E)  $\bowtie_{\text{DNO=Dnumber}}$  DEPARTMENT (D)

## 5) Partition-hash join

- $C_{J4} = 3 * (b_E + b_D) + ((js * r_E * r_D) / bfr_{ED})$   
 $= 3 * (13 + 2,000) + (((1/125) * 10,000 * 125) / 4) = 8,539$   
**Lowest cost!!!**

# Multi-way Join Queries & Join Tree


- Multi-way join queries:  $R1 \bowtie R2 \bowtie \dots \bowtie Rn$ 
  - **TOO MANY** join orders are possible on  $N$ -way joins.
- Join tree:



- Several types: bushy, linear, right-deep, and left-deep
- A (query) plan involving join = a join tree
- A partial plan = a subtree of a join tree

# Types of Join Trees


- **Bushy** (무성한):






# Types of Join Trees (Cont'd)

- *Linear*:




# Types of Join Trees (Cont'd)

- *Right-deep*:



# Types of Join Trees (Cont'd)

- **Left-deep**: gives the following advantages
  - Works well existing algorithms like NLJ or hash-join
  - Facilitates pipelining
  - Dynamic programming can be used with all left-deep trees
    - Optimal solution structure is developed
    - Value of the optimal solution is recursively defined
    - Optimal solution is computed and its value developed in a bottom-up fashion



# Types of Join Trees (Cont'd)

- Number of permutations of left-deep and bushy join trees of  $N$  relations

No. of Relations $N$	No. of Left-Deep Trees $N!$	No. of Bushy Shapes $S(N)$	No. of Bushy Trees $(2N - 2)! / (N - 1)!$
2	2	1	2
3	6	2	12
4	24	5	120
5	120	14	1,680
6	720	42	30,240
7	5,040	132	665,280

# EXAMPLE TO ILLUSTRATE COST-BASED QUERY OPTIMIZATION

---

Chapter 19.6

# Consider the Following Query: Q2

```
SELECT P.Pnumber, P.Dnum, E.Lname, E.Address, E.Bdate  
FROM PROJECT P, DEPARTMENT D, EMPLOYEE E  
WHERE P.Dnum=D.Dnumber AND D.Mgr_ssn=E.Ssn AND P.Plocation='Stafford'
```

- Assume optimizer considers only “left-deep” trees.
- Possible join orderings?
  - 1) ((PROJECT ⋈ DEPARTMENT) ⋈ EMPLOYEE)
  - 2) ((DEPARTMENT ⋈ PROJECT) ⋈ EMPLOYEE)
  - 3) ((DEPARTMENT ⋈ EMPLOYEE) ⋈ PROJECT)
  - 4) ((EMPLOYEE ⋈ DEPARTMENT) ⋈ PROJECT)
- Now evaluate each of the potential join orders based on join cost

# Sample Statistical Information for Relations in Q2

(a): Column information

(b): Table information

(c): Index information

(a)

Table_name	Column_name	Num_distinct	Low_value	High_value
PROJECT	Plocation	200	1	200
PROJECT	Pnumber	2000	1	2000
PROJECT	Dnum	50	1	50
DEPARTMENT	Dnumber	50	1	50
DEPARTMENT	Mgr_ssn	50	1	50
EMPLOYEE	Ssn	10000	1	10000
EMPLOYEE	Dno	50	1	50
EMPLOYEE	Salary	500	1	500

(b)

Table_name	Num_rows	Blocks
PROJECT	2000	100
DEPARTMENT	50	5
EMPLOYEE	10000	2000

(c)

Index_name	Uniqueness	Blevel*	Leaf_blocks	Distinct_keys
PROJ_PLOC	NONUNIQUE	1	4	200
EMP_SSN	UNIQUE	1	50	10000
EMP_SAL	NONUNIQUE	1	50	500

\*Blevel is the number of levels without the leaf level.

# Summary

- Query trees
- Heuristic approaches used to improve efficiency of query execution
- Reorganization of query trees
- Cost-based optimization approach



# References

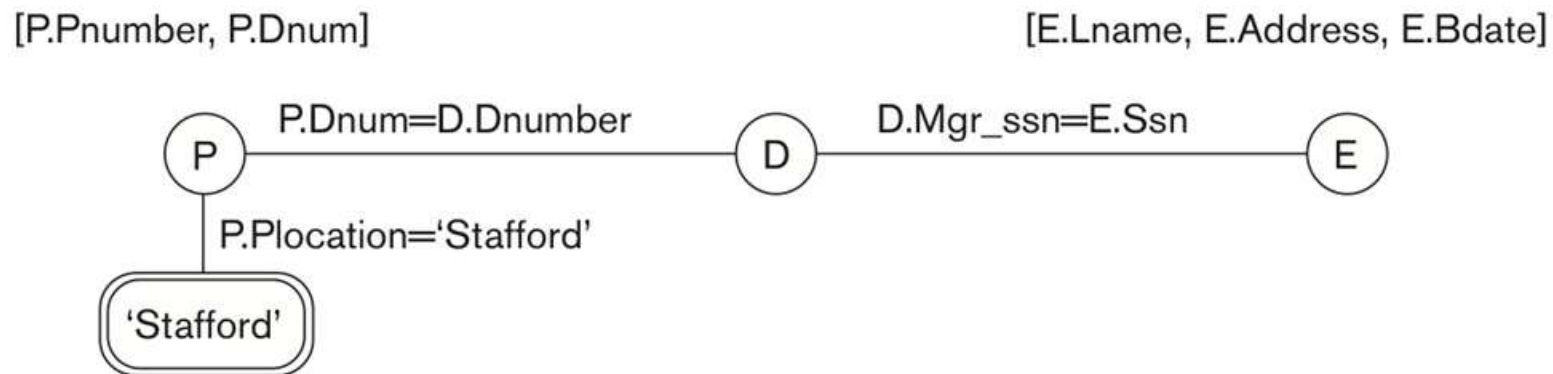
- <https://courses.cs.washington.edu/courses/cse444/14sp/lectures/lecture11-optimization-part2.pdf>

# APPENDIX

---

## [Appendix] Example of a Query Graph

- Q1 can be represented by a *query graph* below:



- Represents *relational calculus expression*
- Relation nodes displayed as *single circles*
- Constants represented by constant nodes: *double circles* (or ovals)
- Selection or join conditions as *edges*
- Attributes to be retrieved displayed in *square brackets*

## Some Transformation Rules

... proving them: ... rules that are useful in query

- Cascade of  $\sigma$ .** A conjunctive selection condition can be broken up into a cascade (that is, a sequence) of individual  $\sigma$  operations:  
 $\sigma_{c_1 \text{ AND } c_2 \text{ AND } \dots \text{ AND } c_n}(R) \equiv \sigma_{c_1}(\sigma_{c_2}(\dots(\sigma_{c_n}(R))\dots))$
- Commutativity of  $\sigma$ .** The  $\sigma$  operation is commutative:  
 $\sigma_{c_1}(\sigma_{c_2}(R)) \equiv \sigma_{c_2}(\sigma_{c_1}(R))$
- Cascade of  $\pi$ .** In a cascade (sequence) of  $\pi$  operations, all but the last one can be ignored:  
 $\pi_{\text{List}_1}(\pi_{\text{List}_2}(\dots(\pi_{\text{List}_n}(R))\dots)) \equiv \pi_{\text{List}_1}(R)$
- Commuting  $\sigma$  with  $\pi$ .** If the selection condition  $c$  involves only those attributes  $A_1, \dots, A_n$  in the projection list, the two operations can be commuted:  
 $\pi_{A_1, A_2, \dots, A_n}(\sigma_c(R)) \equiv \sigma_c(\pi_{A_1, A_2, \dots, A_n}(R))$
- Commutativity of  $\bowtie$  (and  $\times$ ).** The join operation is commutative, as is the  $\times$  operation:  
 $R \bowtie_c S \equiv S \bowtie_c R$   
 $R \times S \equiv S \times R$

Notice that although the order of attributes may not be the same in the relations resulting from the two joins (or two Cartesian products), the *meaning* is the same because the order of attributes is not important in the alternative definition of relation.

- Commuting  $\sigma$  with  $\bowtie$  (or  $\times$ ).** If all the attributes in the selection condition  $c$  involve only the attributes of one of the relations being joined—say,  $R$ —the two operations can be commuted as follows:  
 $\sigma_c(R \bowtie S) \equiv (\sigma_c(R)) \bowtie S$   
 Alternatively, if the selection condition  $c$  can be written as  $(c_1 \text{ AND } c_2)$ , where condition  $c_1$  involves only the attributes of  $R$  and condition  $c_2$  involves only the attributes of  $S$ , the operations commute as follows:  
 $\sigma_c(R \bowtie S) \equiv (\sigma_{c_1}(R)) \bowtie (\sigma_{c_2}(S))$   
 The same rules apply if the  $\bowtie$  is replaced by a  $\times$  operation.
- Commuting  $\pi$  with  $\bowtie$  (or  $\times$ ).** Suppose that the projection list is  $L = \{A_1, \dots, A_n, B_1, \dots, B_m\}$ , where  $A_1, \dots, A_n$  are attributes of  $R$  and  $B_1, \dots, B_m$  are attributes of  $S$ . If the join condition  $c$  involves only attributes in  $L$ , the two operations can be commuted as follows:  
 $\pi_L(R \bowtie_c S) \equiv (\pi_{A_1, \dots, A_n}(R)) \bowtie_c (\pi_{B_1, \dots, B_m}(S))$   
 If the join condition  $c$  contains additional attributes not in  $L$ , these must be added to the projection list, and a final  $\pi$  operation is needed. For example, if attributes

$A_{n+1}, \dots, A_{n+k}$  of  $R$  and  $B_{m+1}, \dots, B_{m+p}$  of  $S$  are involved in the join condition  $c$  but are not in the projection list  $L$ , the operations commute as follows:  
 $\pi_L(R \bowtie_c S) \equiv \pi_L((\pi_{A_1, \dots, A_n, A_{n+1}, \dots, A_{n+k}}(R)) \bowtie_c (\pi_{B_1, \dots, B_m, B_{m+1}, \dots, B_{m+p}}(S)))$   
 For  $\times$ , there is no condition  $c$ , so the first transformation rule always applies by replacing  $\bowtie_c$  with  $\times$ .

- Commutativity of set operations.** The set operations  $\cup$  and  $\cap$  are commutative, but  $-$  is not.
- Associativity of  $\bowtie$ ,  $\times$ ,  $\cup$ , and  $\cap$ .** These four operations are individually associative; that is, if both occurrences of  $\theta$  stand for the same operation that is any one of these four operations (throughout the expression), we have:  
 $(R \theta S) \theta T \equiv R \theta (S \theta T)$
- Commuting  $\sigma$  with set operations.** The  $\sigma$  operation commutes with  $\cup$ ,  $\cap$ , and  $-$ . If  $\theta$  stands for any one of these three operations (throughout the expression), we have:  
 $\sigma_c(R \theta S) \equiv (\sigma_c(R)) \theta (\sigma_c(S))$
- The  $\pi$  operation commutes with  $\cup$ .**  
 $\pi_L(R \cup S) \equiv (\pi_L(R)) \cup (\pi_L(S))$
- Converting a  $(\sigma, \times)$  sequence into  $\bowtie$ .** If the condition  $c$  of a  $\sigma$  that follows a  $\times$  corresponds to a join condition, convert the  $(\sigma, \times)$  sequence into a  $\bowtie$  as follows:  
 $(\sigma_c(R \times S)) \equiv (R \bowtie_c S)$
- Pushing  $\sigma$  in conjunction with set difference.**  
 $\sigma_c(R - S) = \sigma_c(R) - \sigma_c(S)$   
 However,  $\sigma$  may be applied to only one relation:  
 $\sigma_c(R - S) = \sigma_c(R) - S$
- Pushing  $\sigma$  to only one argument in  $\cap$ .**  
 If in the condition  $\sigma_c$  all attributes are from relation  $R$ , then:  
 $\sigma_c(R \cap S) = \sigma_c(R) \cap S$
- Some trivial transformations.**  
 If  $S$  is empty, then  $R \cup S = R$   
 If the condition  $c$  in  $\sigma_c$  is true for the entire  $R$ , then  $\sigma_c(R) = R$ .

NOT  $(c_1 \text{ AND } c_2) \equiv (\text{NOT } c_1) \text{ OR } (\text{NOT } c_2)$   
 NOT  $(c_1 \text{ OR } c_2) \equiv (\text{NOT } c_1) \text{ AND } (\text{NOT } c_2)$

# Example Table for Calculating Execution Cost

Cost Estimate Items	Number
$r_{E(MPLOYEE)}$	10,000 records
$b_{E(MPLOYEE)}$	2,000 blocks
$r_{D(DEPARTMENT)}$	125 records
$b_{D(DEPARTMENT)}$	13 blocks
$x_{Dnumber}$ (Number of level of a primary index on the $Dnumber$ attribute in $DEPARTMENT$ )	1
$x_{Mgr\_ssn}$ (Number of level of a secondary index on the $Mgr\_ssn$ attribute in $DEPARTMENT$ )	2
$s_{Mgr\_ssn}$ (Selectivity cardinality on $Mgr\_ssn$ )	1
Join selectivity ( $js$ ) on $EMPLOYEE \bowtie_{Dno=Dnumber} DEPARTMENT$	$\frac{1}{\max(NDV(Dno, EMPLOYEE), NDV(Dnumber, DEPARTMENT))} = \frac{1}{125}$
Blocking factor of the resulting join ( $bfr_{ED}$ )	4 records

$DEPARTMENT.Dnumber$ : a key attribute of  $DEPARTMENT$