



TRANSACTION PROCESSING II

Chapter 21

- Concurrency control

Prof. Young-Kyoon Suh

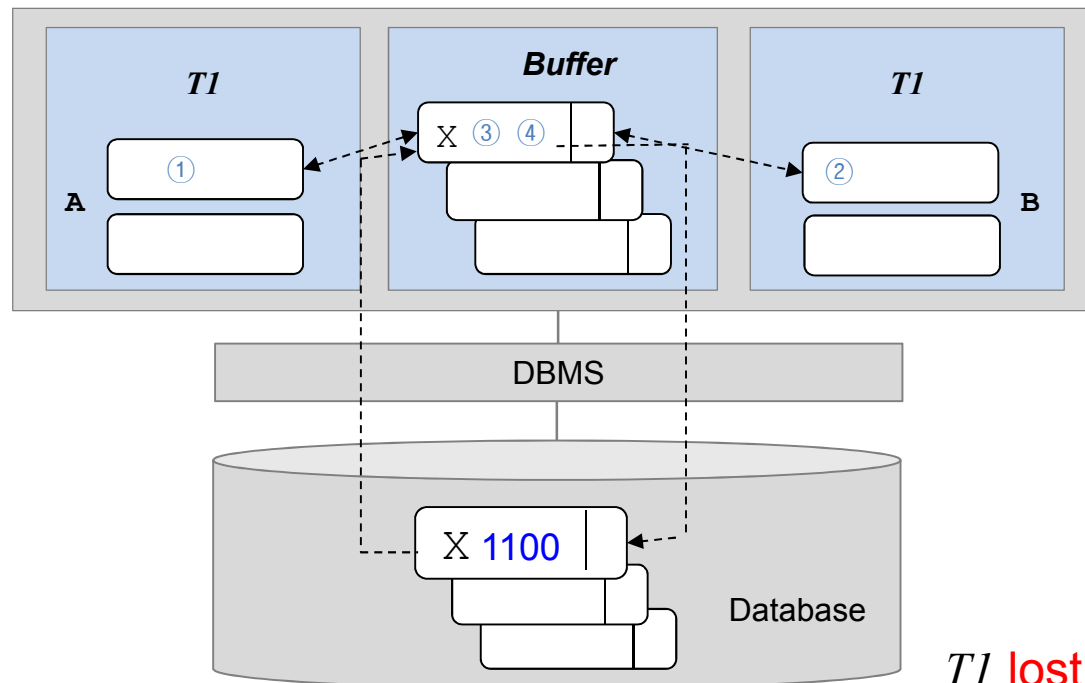
Introduction

- Concept of *Concurrency Control*
 - If concurrent transactions have access to the same data item, concurrency control techniques effectively orchestrates the access control of the item, not to harm consistency of the item.
 - Defines a set of rules to guarantee serializability
- Scenarios of read/write from two concurrent transactions:
T1 and *T2*

	<i>T1</i>	<i>T2</i>	Problem	Concurrency
Scenario 1	Read	Read	No	Allowed
Scenario 2	Read	Write	<i>Dirty read,</i> <i>Non-repeatable read,</i> <i>Phantom read</i>	Allowed or Disallowed along with isolation level
Scenario 3	Write	Write	<i>Lost Update</i> (never allowed)	Disallowed; but possible with locks

Introduction – Lost Update (Cont'd)

$T1$	$T2$	Buffer Data
$A = \text{read_item}(X);$ ① $A = A - 100;$		$X = 1000$
	$B = \text{read_item}(X);$ ② $B = B + 100;$	$X = 1000;$
$\text{write_item}(A \rightarrow X);$ ③		$X = \mathbf{900}$
	$\text{write_item}(B \rightarrow X);$ ④	$X = \mathbf{1100}$



$T1$ **lost** its update by $T2$.

Introduction (Cont'd)

- To solve the *lost update* problem, use **locking** techniques.
 - **Lock**: a *variable* associated with a data item
 - Describes the status of the item, with respect to possible operations that can be applied to it
 - E.g., “My item is locked. You cannot use mine right now.”
 - Used as a means of **synchronizing** the access by concurrent transactions to the database items.
 - One lock for each item in the database

T1

T2

Time

```
SQL> SET TRANSACTION NAME 'T1';
```

```
Transaction set.
```

```
SQL>
```

```
SQL> SELECT *
2 FROM BOOK
3 WHERE bookid = 1;
```

BOOKID	BOOKNAME	PUBLISHER	PRICE
1	Database	Pearson	20000

```
SQL> UPDATE BOOK
```

```
2 SET price = 21000
3 WHERE bookid = 1;
```

```
1 row updated.
```

Now the Book tuple is **locked** by T1.

```
SQL> _
```

```
SQL> SET TRANSACTION NAME 'T2';
```

```
Transaction set.
```

```
SQL>
```

```
SQL> SELECT *
2 FROM BOOK
3 WHERE bookid = 1;
```

BOOKID	BOOKNAME	PUBLISHER	PRICE
1	Database	Pearson	20000

```
SQL> UPDATE BOOK
```

```
2 SET price = 21000
3 WHERE bookid = 1;
```

```
_
```

T2 waits until the **lock** on the Book tuple is released.

```
SQL> SELECT *
2 FROM BOOK
3 WHERE bookid = 1;
```

BOOKID	BOOKNAME	PUBLISHER	PRICE
1	Database	Pearson	21000

```
SQL>
```

```
SQL> commit;
```

```
Commit complete.
```

```
SQL>
```

```
1 row updated.
```

```
SQL>
```

TWO-PHASE LOCKING TECHNIQUES FOR CONCURRENCY CONTROL

Chapter 20.1

Types of Locks and System Lock Tables

1) *Binary locks* (used by the binary locking scheme)

- Two states (values)
 - *Locked* (1): Item cannot be accessed
 - *Unlocked* (0): Item can be accessed when requested
- Locking notation: **lock (X)**
 - Indicates the current status (value) of the lock associated with item X
- Two operations are used with binary locks:
 - **lock_item**
 - Must be invoked *before* use
 - **unlock_item**
 - Must be invoked *after* use

```
lock_item(X):  
B:  if LOCK(X) = 0          (*item is unlocked*)  
    then LOCK(X) ← 1      (*lock the item*)  
    else  
        begin  
            wait (until LOCK(X) = 0  
                and the lock manager wakes up the transaction);  
            go to B  
        end;  
unlock_item(X):  
    LOCK(X) ← 0;          (* unlock the item *)  
    if any transactions are waiting  
        then wakeup one of the waiting transactions;
```

Types of Locks and System Lock Tables (Cont'd)

1) Binary locks (Cont'd)

- Easy to implement via a binary-valued variable, say **LOCK**
- Each lock can be a record with three fields:
<Data_item_name, LOCK, Locking_Transaction>.
- **Lock table**: implemented as a **hash file** on the item name
 - Used to maintain these records only for the items currently locked
 - Specifies data items that have locks
- **Lock manager subsystem** of the DBMS
 - Keeps track of and controls access to locks
 - Enforces locking rules (“call `lock_item` before use”).
- At most one transaction can hold the lock on an item at a given time
- Binary locking **too restrictive** for database items

Types of Locks and System Lock Tables (Cont'd)

2) **Shared/Exclusive** (or **Read/Write**) Locks (used by the shared/exclusive locking scheme)

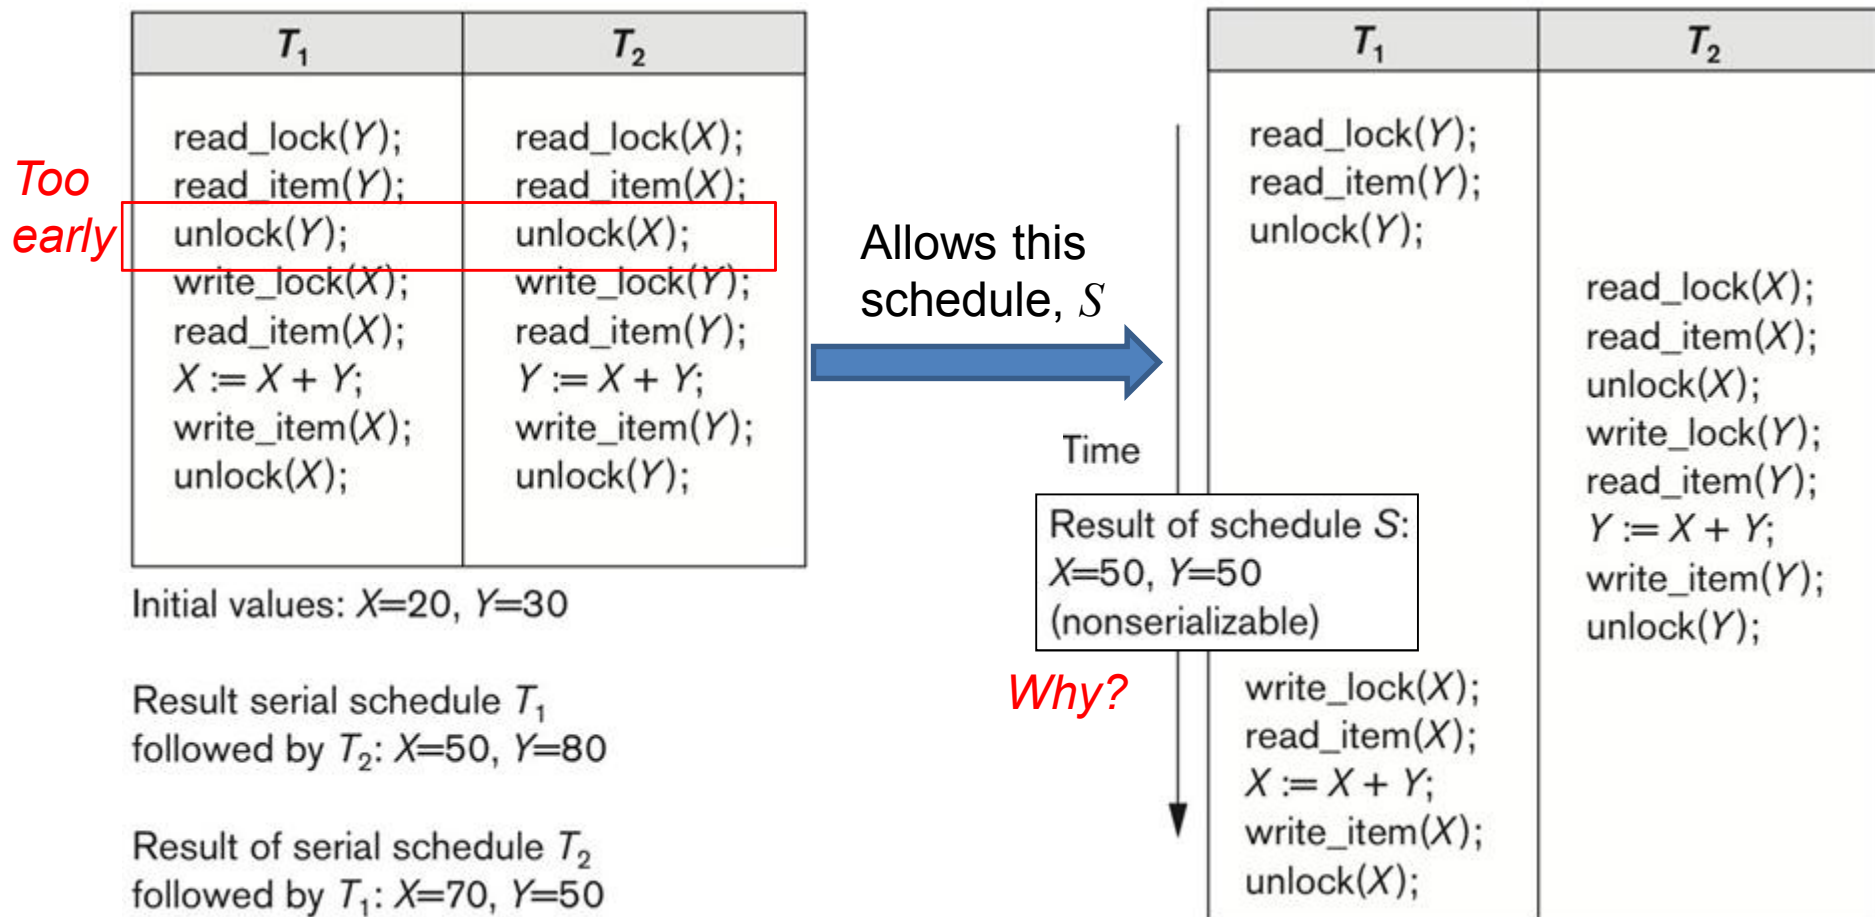
- Called a *multiple-mode* lock
- Three states (values) of `lock(X)`
 - *read-locked* (or shared-locked): Item can be read by other transactions.
 - *write-locked* (or exclusive-locked): A single transaction *exclusively* holds the lock on the item; Item can **NOT** be accessed for a shared-lock request.
 - *unlocked*: Item can be accessed
- Has three locking operations:
 - `read_lock(X)`, `write_lock(X)`: Must be invoked *before* use
 - `unlock(X)`: Must be invoked *after* use
- Each record in the lock table has four fields:
 - `<Data_item_name, (read|write) LOCK, # of reads, Locking_Transaction(s)>`.
- System maintains lock records *only for locked items* in the table.
 - If an item not locked, it would not be in the table.

Conversion of Locks

- Transaction that already holds a lock is allowed to “convert” the lock from one state to another.
 - Called *lock conversion*
- *Upgrading*
 - Issues `read_lock(X)` followed by `write_lock(X)`
- *Downgrading*
 - Issues `write_lock(X)` followed by `read_lock(X)`

Problems of the Previous Locking Schemes

- Using binary locks or read/write locks in transactions *doesn't guarantee serializability of schedules* on its own.



TWO-PHASE LOCKING

- Guarantees Serializability (as well as solves the lost update problem)

Two-Phase Locking Protocol

- All locking operations precede the *first* unlock operation in the transaction: lock requests -> lock releases (**never backward**).
- Consists of two phases:
 - *Expanding* (growing) phase
 - New locks can be acquired but none can be released
 - Don't release acquired locks
 - Lock conversion *upgrades* must be done during this phase.
 - *Shrinking* phase
 - Existing locks can be released **BUT** none can be acquired.
 - Once you begin to unlock, then you can't acquire a lock until all the locks are released.
 - *Downgrades* must be done during this phase.

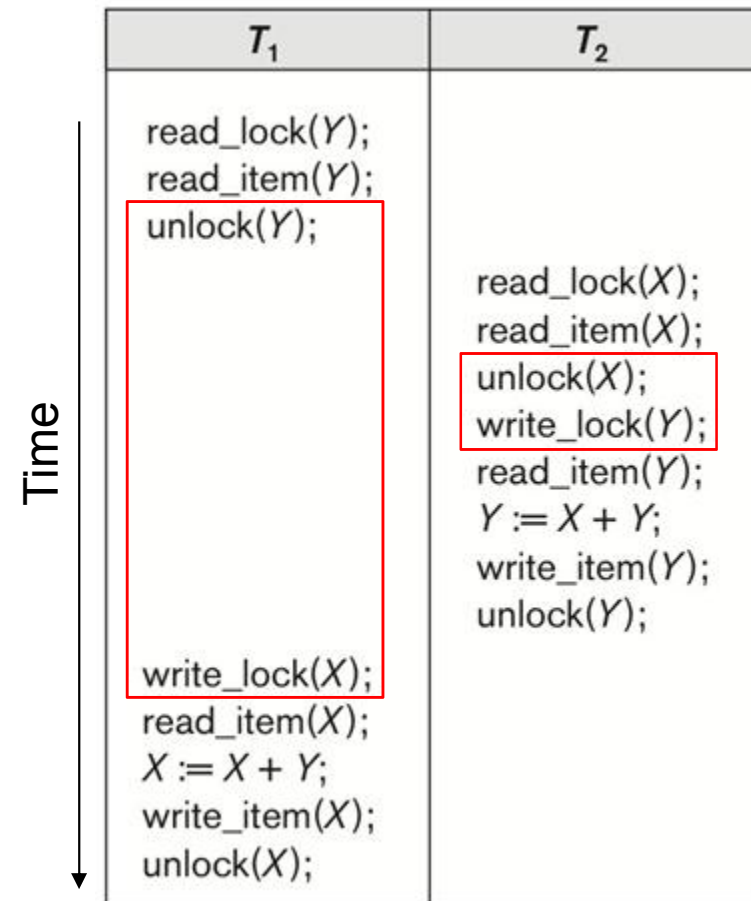
Review: Transactions That Don't Obey 2PL

T_1	T_2
<pre> read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X); </pre>	<pre> read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y); </pre>

Initial values: $X=20$, $Y=30$

Result serial schedule T_1
followed by T_2 : $X=50, Y=80$

Result of serial schedule T_2
followed by T_1 : $X=70, Y=50$



Result of schedule S:
X=50, Y=50
(nonserializable)

Transactions Following 2PL

But **deadlock** happens

T_1'	T_2'
read_lock(Y); read_item(Y); write_lock(X); unlock(Y) read_item(X); $X := X + Y$; write_item(X); unlock(X);	read_lock(X); read_item(X); write_lock(Y); unlock(X) read_item(Y); $Y := X + Y$; write_item(Y); unlock(Y);

2PL guarantees serializability of these two transactions!

Variations of 2PL

- Basic 2PL: covered
- *Conservative* 2PL
 - Requires a transaction to lock “all” the items it accesses *before the transaction begins execution*
 - How? By predeclaring (i) **read-set** and (ii) **write-set**
 - Waits until needed items are available for locking, if *any* of the predeclared items cannot be locked.
 - Once the transaction starts, it’s already in the **shrinking** phase.
 - *Different* from rigorous 2PL to be discussed soon shortly
 - **Deadlock-free** protocol; but not in practice; Why?

Variations of 2PL – Strict/Rigorous 2PL (Cont'd)

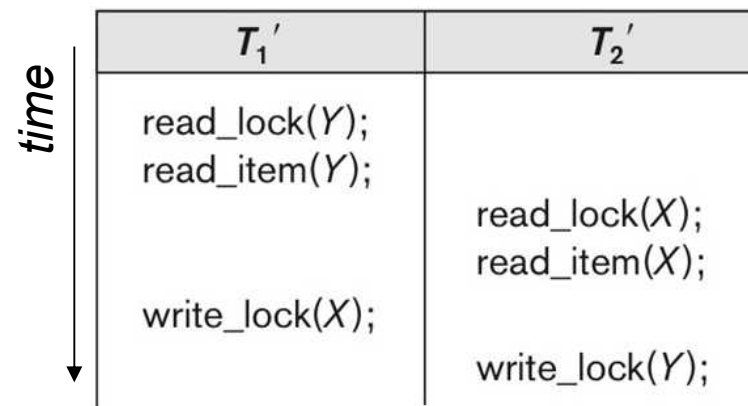
	<i>Strict</i> 2PL	<i>Rigorous</i> 2PL
Common	Guarantees <i>strict</i> schedules (in terms of recoverability)	
	NOT deadlock-free	
Differences	<i>T</i> does <i>not</i> releases <i>write</i> locks <i>until after T</i> commits or aborts.	<i>T</i> does <i>not</i> releases <i>any</i> (read or write) locks <i>until after T</i> commits or aborts.
	<i>No other transaction can read/write an item written by T unless T has committed</i>	
	<i>Harder</i> to implement than rigorous 2PL	<i>Easier</i> to implement than rigorous 2PL
	-	In the expanding phase until transaction ends - <i>Different from conservative 2PL</i>

HOW TO RESOLVE DEADLOCK?

- Deadlock prevention
- Deadlock detection

Dealing with *Deadlock*

- Locking reveals three major problems:
 - 1) A **high** overhead (lock request followed by read or write),
 - 2) **Deadlock**, and 3) **Starvation**
- **Deadlock**:
 - Occurs when each transaction T in a set is waiting for some item locked by some other transaction T' in the set.
 - Both transactions are **stuck** in a waiting queue



[A partial schedule of T_1' and T_2']

Dealing with Deadlock: *Deadlock Prevention* Protocols

1) *Conservative 2PL*: deadlock-free

- Impractical; Limits concurrency

2) *Another protocol*:

- *Orders all the items* in the database, and makes sure that a transaction that needs several items will lock them according to that order.
 - Limits concurrency as well
 - The programmer (or the system) is aware of the chosen order of the items, which is also impractical

3) *Timestamp(TS)-based protocol*:

- Suppose that for 2 transactions (T_1, T_2), T_1 tries to lock an item locked by T_2 .
- *Wait-die* (scheme): **non-preemptive** from the perspective of lock requestor
 - If $TS(T_1) < TS(T_2)$, then T_1 waits; otherwise, abort and restart T_1 with same TS.
- *Wound-wait* (scheme): **preemptive** from the perspective of lock requestor
 - If $TS(T_1) < TS(T_2)$, then aborts and restart T_2 with same TS; otherwise, T_1 waits.

Dealing with Deadlock: *Deadlock Prevention* Protocols (Cont'd)

4) *No waiting algorithm* (not requiring TS)


- If transaction is unable to obtain a lock, then it's immediately aborted and then restarted *after a delay*

5) *Cautious waiting algorithm: deadlock-free*

- Suppose that T_1 tries to lock an item locked by T_2 .
- If T_2 is not blocked (or, not waiting for some other locked item), then T_1 is *blocked and allowed to wait*; otherwise, T_1 is aborted.
- Designed to reduce # of needless restarts/aborts

Dealing with Deadlock: *Deadlock Detection*

- The DBMS checks if a state of deadlocks actually exists.
- A simple way: drawing a *wait-for* graph
 - Lock wait -> add an edge,
 - Lock release -> drop that edge
 - If a cycle exists, then that's the state of **deadlock**.
 - Checking for a cycle everytime: *too much overhead*
 - One solution: doing the cycle check periodically, based on (i) *# executing transactions counts* or (ii) *period of time waiting for locking items*
 - *Victim selection*
 - Chooses which transaction to *abort*
 - Criteria: younger transactions that have not made many changes; Why so ?



[A wait-for graph for the partial schedule]

Dealing with Deadlock: *Timeouts*

- Another simple scheme for deadlock detection
 - Low overhead and simplicity
 - If T times out beyond a system-defined threshold, then T aborted
 - Regardless of whether the deadlock actually exists

Dealing with Deadlock: *Starvation*

- Occurs when a transaction cannot proceed for infinite time
 - Also occurs for victim selection on the same transaction
 - Solution 1: *first-come, first-served*
 - Called a *fair-waiting scheme*
 - Solution 2
 - Giving some priority to transactions so that they can eventually proceed
 - Other solutions:
 - Wait-die, or wound-wait

Summary

- Concept of Concurrency Control
- Lock Types and Lock Conversion
- Two-Phase Locking
- Variation of Two-Phase Locking Schemes
- Deadlock Detection, Prevention, Starvation

APPENDIX

Types of Locks and System Lock Tables (Cont'd)

1) Binary locks (Cont'd)

- Rules enforced by the binary locking scheme
 - A transaction (T) must issue `lock_item(X)` before any `read_item(X)` or `write_item(X)`.
 - T must issue `unlock_item(X)` after all `read_item(X)` and `write_item(X)` are completed.
 - T will not issue `lock_item(X)` if it already holds the lock on item X .
 - T will not issue `unlock(X)` unless it already holds the lock on item X .

Types of Locks and System Lock Tables (Cont'd)

2) *Shared/Exclusive* (or *Read/Write*) Locks (Cont'd)

- Locking and unlocking operations for two-mode (read/write, or shared/exclusive) locks

```
read_lock(X):
B:  if LOCK(X) = "unlocked"
    then begin LOCK(X) ← "read-locked";
        no_of_reads(X) ← 1
    end
    else if LOCK(X) = "read-locked"
        then no_of_reads(X) ← no_of_reads(X) + 1
    else begin
        wait (until LOCK(X) = "unlocked"
            and the lock manager wakes up the transaction);
        go to B
    end;

write_lock(X):
B:  if LOCK(X) = "unlocked"
    then LOCK(X) ← "write-locked"
    else begin
        wait (until LOCK(X) = "unlocked"
            and the lock manager wakes up the transaction);
        go to B
    end;

unlock (X):
    if LOCK(X) = "write-locked"
        then begin LOCK(X) ← "unlocked";
            wakeup one of the waiting transactions, if any
        end
    else if LOCK(X) = "read-locked"
        then begin
            no_of_reads(X) ← no_of_reads(X) - 1;
            if no_of_reads(X) = 0
                then begin LOCK(X) = "unlocked";
                    wakeup one of the waiting transactions, if any
                end
        end
    end;
```

Types of Locks and System Lock Tables (Cont'd)

2) *Shared/Exclusive* (or *Read/Write*) Locks (Cont'd)

- Rules enforced by the shared/exclusive locking scheme
 - A transaction (T) must issue `read_lock(X)` or `write_lock(X)` before any `read_item(X)`.
 - T must issue `write_lock(X)` before any `write_item(X)`.
 - T must issue `unlock(X)` after all `read_item(X)` and `write_item(X)`.
 - T will not issue `read_lock(X)` if it already holds a read or a write lock on item X .
 - This rule may be relaxed for *downgrading* of locks.
 - T will not issue `write_lock(X)` if it already holds a read or a write lock on item X .
 - This rule may also be relaxed for *upgrading* of locks.
 - T will not issue `unlock(X)` unless it already holds a read lock or a write lock on item X .